

## Structured Query Language

Relational algebra describes operations we can perform on relations that we can use to answer queries. It leads to a very procedural query language: you have to describe the sequence of operations to apply in detail.

A declarative query language only requires you to describe what you want to query, not how.

SQL (Structured Query Language) has become the standard for declarative query languages over relational data. It is one of the reasons for the success of relational database systems. It has been standardized, and you can think of it as a common API for relational databases.

Examples will be based on the following relations:

Account:	account	type	balance
	991	CH	1000
	992	SA	500
	993	CH	50
	994	CH	100
	995	CD	10000

Client:	ssn	name
	1111	R
	2222	T
	3333	A

HasAccount:	account	ssn
	991	1111
	992	1111
	993	2222
	994	1111
	994	3333
	995	1111

**Basic queries:** A basic SQL query looks as follows:

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualifications
```

where *target-list* is a list of attribute names,<sup>1</sup> *relation-list* is a list of table names, and *qualifications* is a list of conditions on tuples, conditions involving comparisons of the form *attribute op attribute* or *attribute op constant* where *op* is a comparison operator such as =, >, <, >=, <=, <>, and combined using AND, OR, and NOT. The DISTINCT qualifier indicates that the tuples in the result should not contain any duplicates.

**Conceptual evaluation strategy:**

1. Compute product of tables in *relation-list*
2. Discard tuples that fail *qualifications*
3. Discard attributes not in *target-list*
4. If DISTINCT, eliminate duplicate tuples

With DISTINCT this roughly translates to the relational algebra expression:

$$\pi_{\text{target-list}}(\sigma_{\text{qualifications}}(r_1 \times \cdots \times r_k))$$

where  $r_1, \dots, r_k$  are the relations in *relation-list*. (Without DISTINCT, we need a version of relational algebra with multisets, as discussed last week.)

The above conceptual evaluation strategy is just that, conceptual. It is not the most efficient way of evaluating queries, but it establishes a reference.

Example:

```
SELECT C.name
FROM Client C, HasAccount H
WHERE C.ssn = H.ssn AND H.account = 994
```

yielding the relation:

C.name
R
A

---

<sup>1</sup>The special symbol \* means *all attributes*.

```

SELECT DISTINCT H.ssn
FROM HasAccount H, Account A
WHERE A.account = H.account AND A.type = 'CH'

```

yielding the relation:

H.ssn
1111
2222
3333

**Nested queries and set conditions:** Here is another query to return the name of clients with account 994, this time using a nested query (which behaves like a set):

```

SELECT C.name
FROM Client C
WHERE C.ssn IN (SELECT H.ssn
                FROM HasAccount H
                WHERE H.account = 994)

```

Conceptually, this checks the qualifications of each tuple by computing the subquery, and checking whether `C.ssn` appears in the result column of the subquery (which is required to be a one-column relation, whose column name is irrelevant). Note that in this case the subquery can be computed once and its result used for every tuple.

```

SELECT C.name
FROM Client C
WHERE EXISTS (SELECT *
              FROM HasAccount H
              WHERE H.ssn = C.ssn and H.account = 994)

```

Conceptually, this checks whether the relation described by the subquery is non-empty. We can check that this computed the same result as above, although note that now the subquery must be explicitly re-evaluated for every tuple of `Client`, since it depends on `C.ssn`.

The following query uses `IN` in a way that is more difficult to avoid (unlike the above examples). It returns the clients with both a checking and a savings account:

```

SELECT H.ssn

```

```

FROM HasAccount H, Account A
WHERE H.account = A.account AND A.type = 'CH'
    AND H.ssn IN (SELECT H2.ssn
                  FROM HasAccount H2, Account A2
                  WHERE H2.account = A2.account AND A2.type = 'SA')

```

**Aggregation:** Lets you aggregate values of attributes across tuples.

Types of aggregation:

```

SUM ([DISTINCT] attr) AS name
AVG ([DISTINCT] attr) AS name
MAX (attr) AS name
MAX (attr) AS name
COUNT ([DISTINCT] attr) AS name

```

Examples:

```

SELECT COUNT(C.ssn) AS count
FROM Client C

```

yielding:

count
3

```

SELECT SUM(A.balance) AS total
FROM Account A

```

yielding:

total
11650

```

SELECT AVG(A.balance) AS avg_checking
FROM Account A
WHERE A.type = 'CH'

```

yielding:

avg_checking
383.33

```
SELECT SUM(A.balance) AS total_1111
FROM HasAccount H, Account A
WHERE H.account = A.account AND H.ssn = 1111
```

yielding:

total_1111
2900

**Aggregation with grouping:** Lets you aggregate values of attributes across tuples, partitioned into groups defined by the value of a set of grouping attributes.

Example:

```
SELECT H.ssn, SUM(A.balance) as total
FROM HasAccount H, Account A
WHERE H.account = A.account
GROUP BY H.ssn
```

yielding:

H.ssn	total
1111	11600
2222	50
3333	100

**Conceptual evaluation strategy:** The general form of a query with aggregation and grouping is:

```
SELECT target-list
FROM relation-list
WHERE qualifications
GROUP BY grouping-list
HAVING group-qualifications
```

1. Compute product of tables in *relation-list*
2. Discard tuples that fail *qualifications*
3. Discard attributes not in *target-list*
4. Partition into groups by values of attributes in *grouping-list*
5. Discard groups that fail *group-qualifications*
6. Aggregate tuples in each group

Example: Maximum balance of clients with more than one non-CD account.

```
SELECT H.ssn, MAX(A.balance) as max
FROM HasAccount H, Account A
WHERE A.type <> 'CD'
GROUP BY H.ssn
HAVING COUNT(A.balance) > 1
```

**SQL and CRUD operations** We can think of a **SELECT** query as a read operation on the tuples in relations. What about the other CRUD operations? SQL provides ways to create, update, and delete tuples from relations:

```
INSERT INTO table
VALUES (value, ..., value)
```

```
UPDATE table
SET attribute = value,
    ...,
    attribute = value
WHERE qualifications
```

```
DELETE FROM table
WHERE qualifications
```

Database systems also provide a Data Definition Language (DDL) which provide CRUD operations at the level of relations themselves: creating (**CREATE TABLE**), updating (**ALTER TABLE**), and deleting relations(**DROP TABLE**) from the database.