

# Exploring the use of B+ trees in Database Management Systems

---

*We're in the endgame now*

# INDEXING

**Indexing** is a key concept in Database Management Systems, as one optimize performance when a query is processed.

The time complexity heavily depends on the **data structure** used.

# UNSORTED ARRAY

10	5	7	18	1	2
----	---	---	----	---	---

**insert:**  $O(1)$

**search:**  $O(N)$


## SORTED ARRAY

1	2	5	7	10	18
---	---	---	---	----	----

**insert:**  $O(N)$

**search:**  $O(\log n)$


# HASHMAP



**insert:**  $\mathcal{O}(1)$

**search:**  $\mathcal{O}(1)$

# BINARY SEARCH TREES



**insert:**  $O(\log N)$

**search:**  $O(\log N)$  /  $O(N)$

---

---

Each time you retrieve a node from a tree you have to pull a block from memory. Since a BST node can only point to 2 children, it's inefficient to retrieve the values stored in the BST as you need to make many memory calls.


**In order to minimize the number of memory calls  
B+ trees are used.**

---

## **B+ Tree Terminology**




# B+ TREES



# INTERNAL NODES


Internal Node

Node Pointer




# INTERNAL NODES

Internal Node



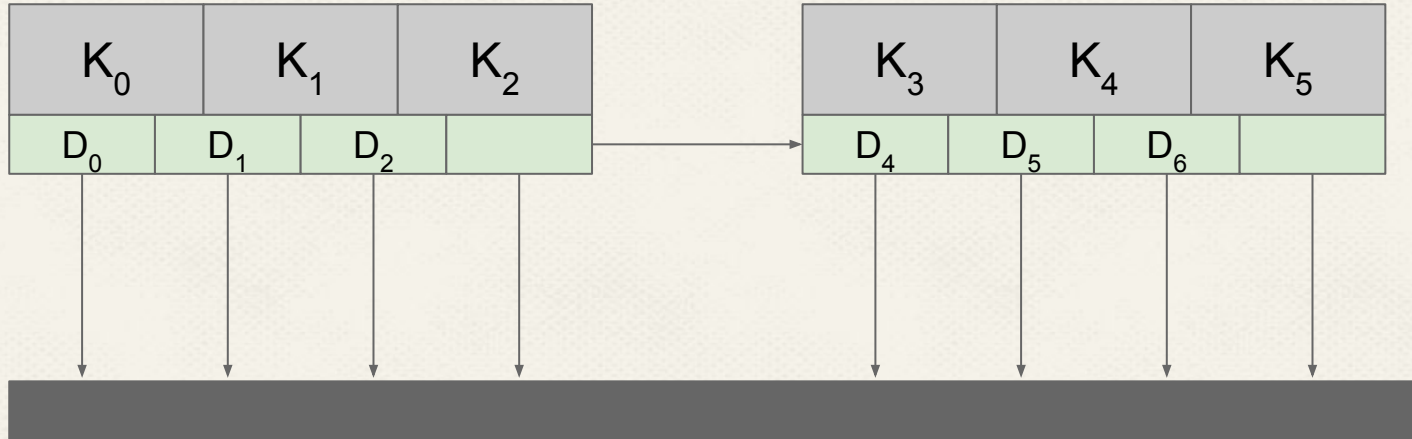
# LEAF NODES

Leaf Node



# LEAF NODES

Leaf Node



## Degree and order

*Order of B+ trees:*

The maximum number of children that a B+ tree node can have.


*Degree of B+ trees:*

For a B+ tree of degree **d**, each node must have **at least d-1 keys** can have **at most 2d-1 keys**.


---


# Example

*B+ tree of degree 2*










Keys  $\geq 4$  and  $< 7$



---

# INSERT


*How do we add a value?*

# Insertion Algorithm

1. Find the node that the key should be inserted in.
2. Check if node is full.
  - a) If node is not full, insert the key in the correct position.
  - b) If the node is full, split the node into two. Copy and insert the median key value into the parent if the node is a leaf node. Else, promote the median without copying


# Insertion Algorithm

## Case - 1: Insert 3 and 5



# Insertion Algorithm


## Case - 1: Insert 3 and 5





# Insertion Algorithm

## Case - 2: Insert 7





Node is full





# Insertion Algorithm

## Case - 2: Insert 7




# Insertion Algorithm

## Case - 2: Insert 7



# Insertion Algorithm

## Case - 2: Insert 7



---

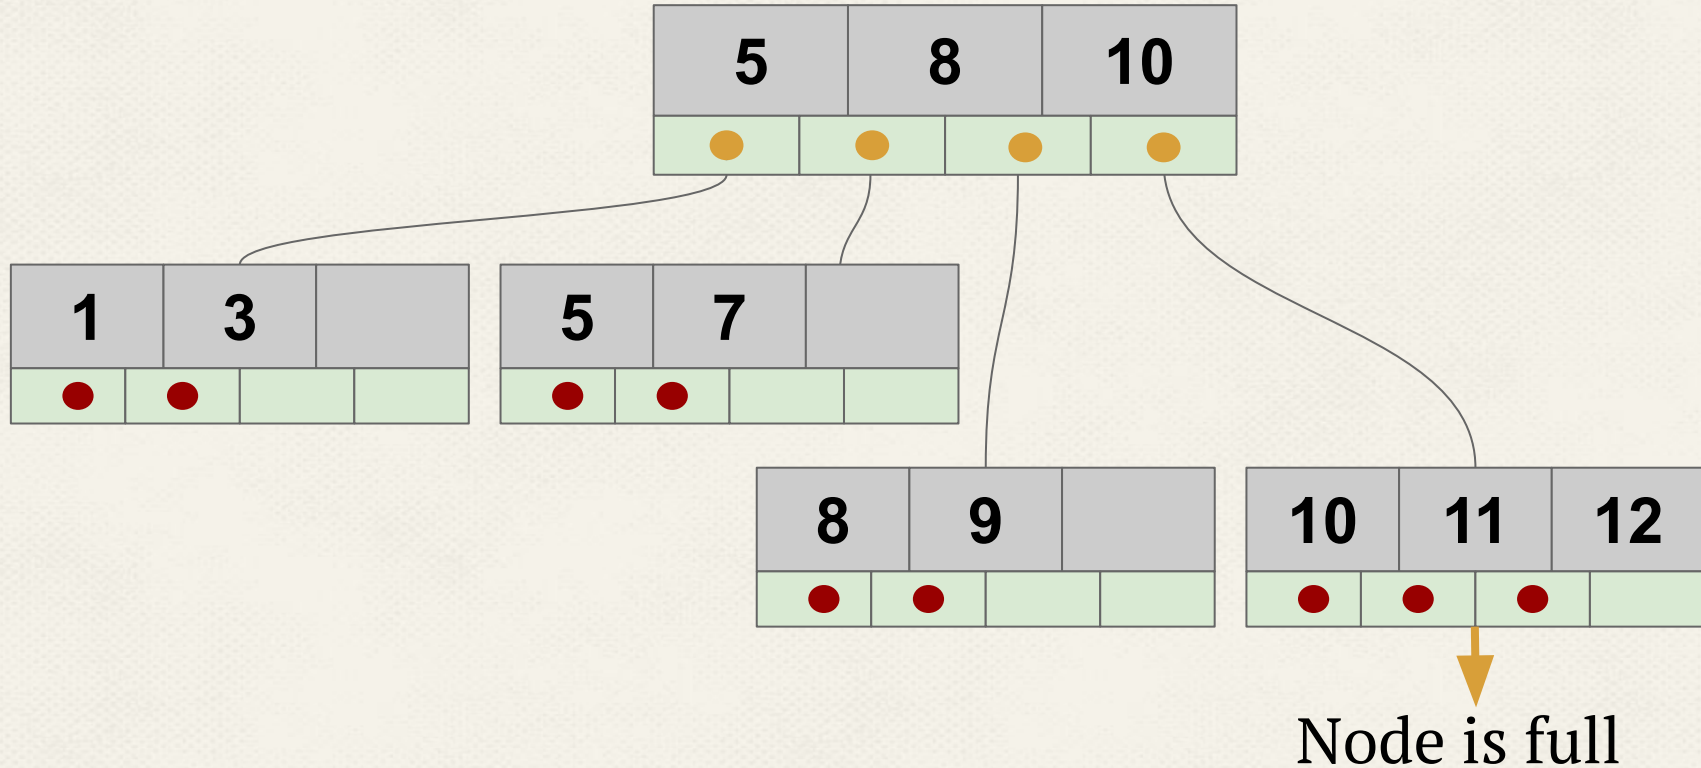
---

**Fast forward**




# Insertion Algorithm

## Case - 3: Insert 13




# Insertion Algorithm

## Case - 3: Insert 13




# Insertion Algorithm

## Case - 3: Insert 13



# Insertion Algorithm

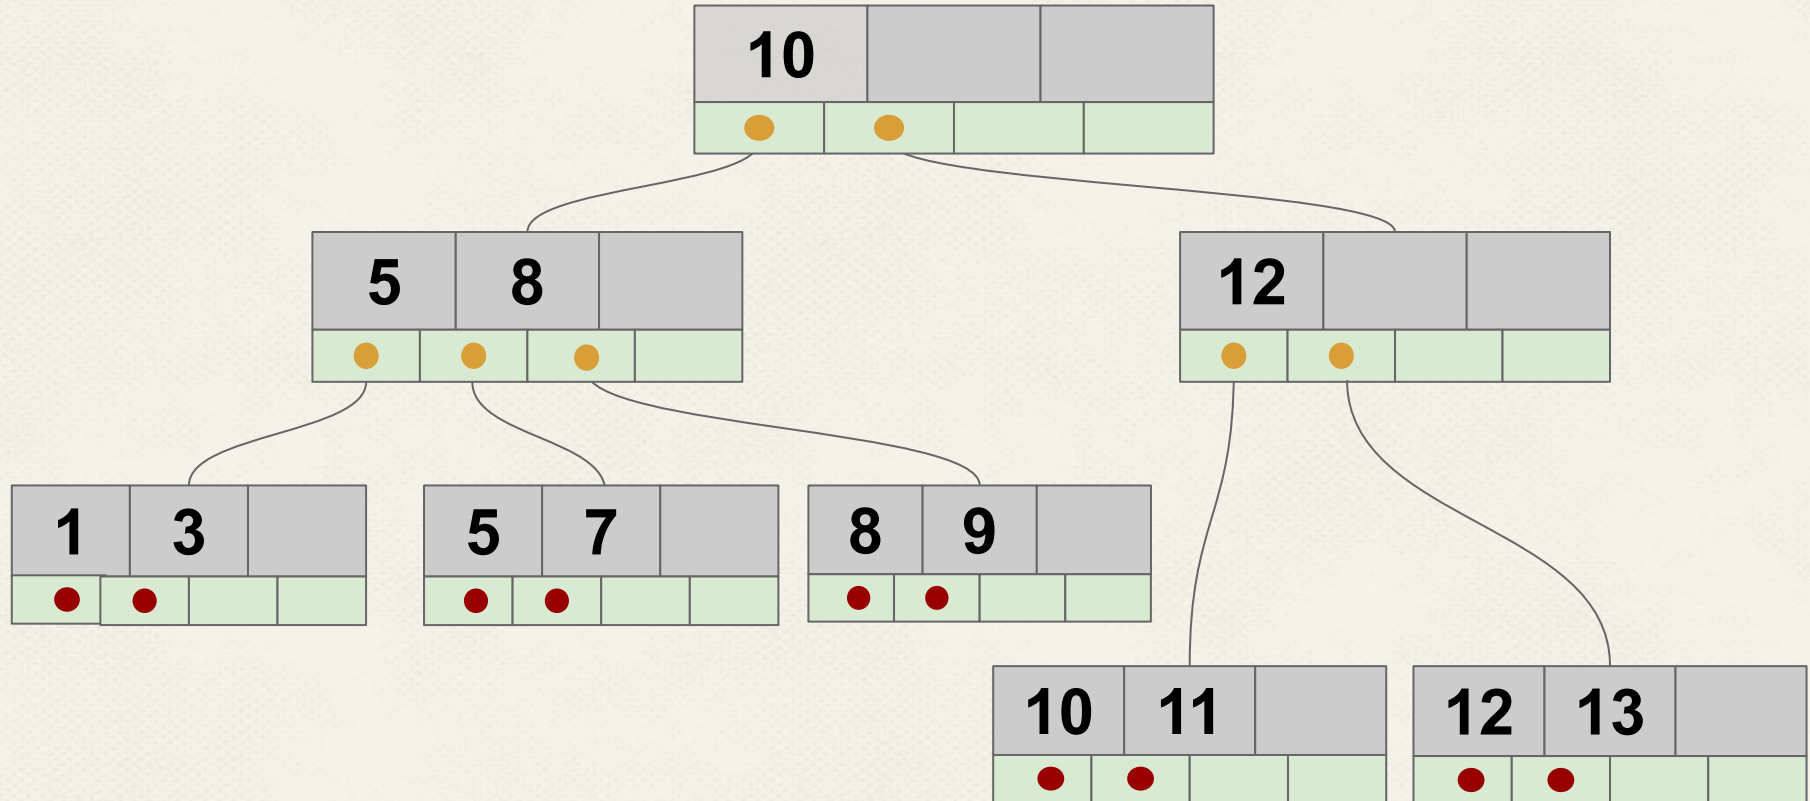
## Case - 3: Insert 13





# Insertion Algorithm

## Case - 3: Insert 13



---

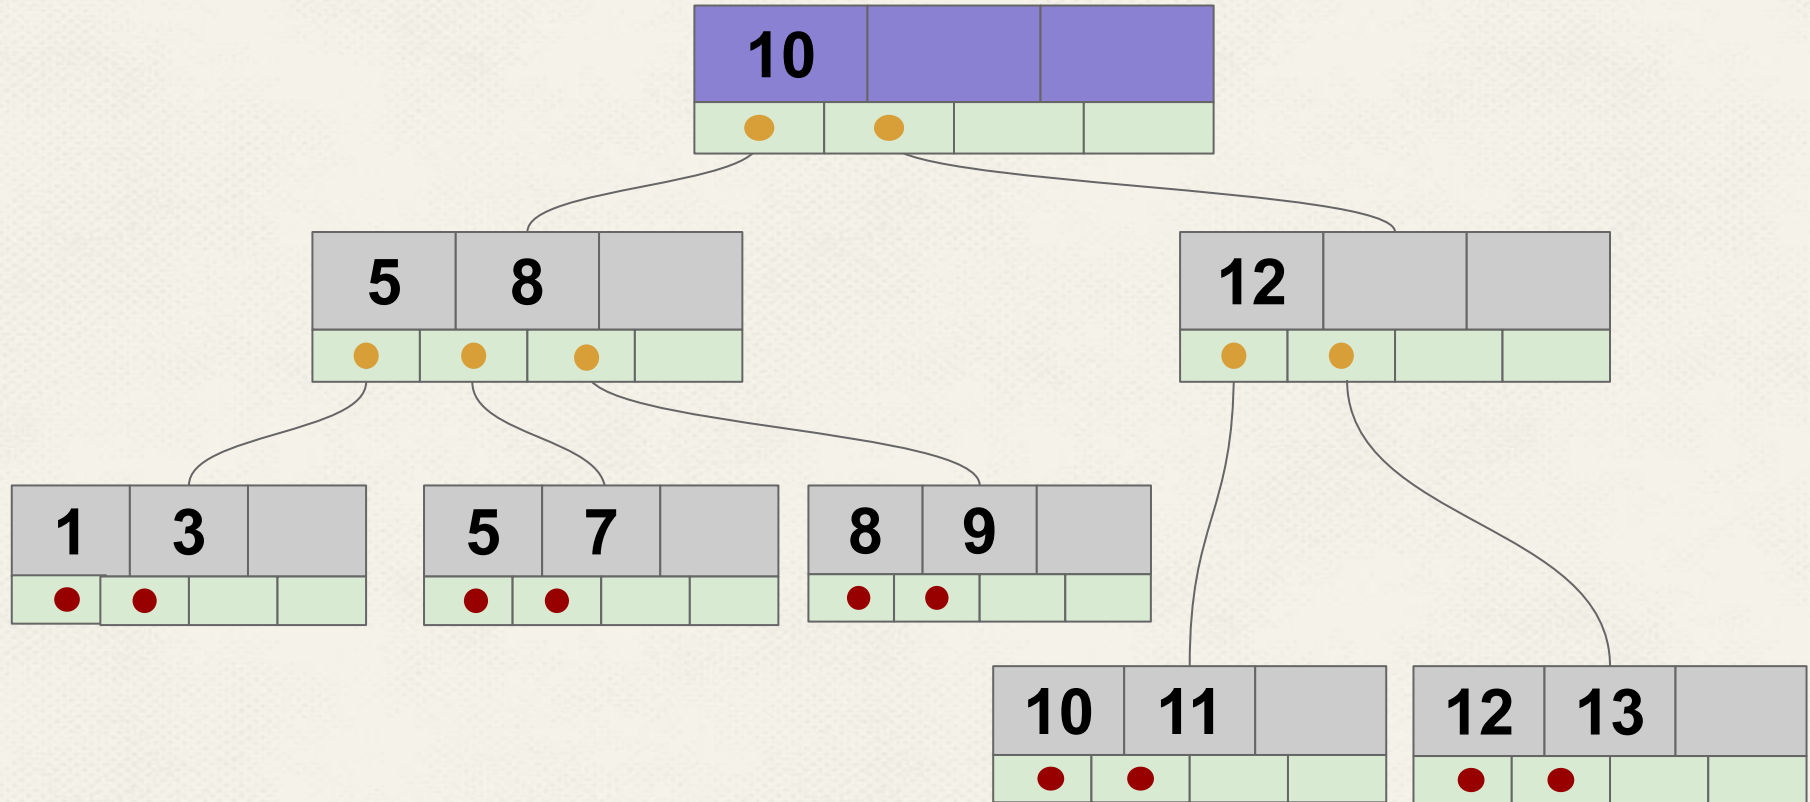
# SEARCH

*How do we find the data we want?*

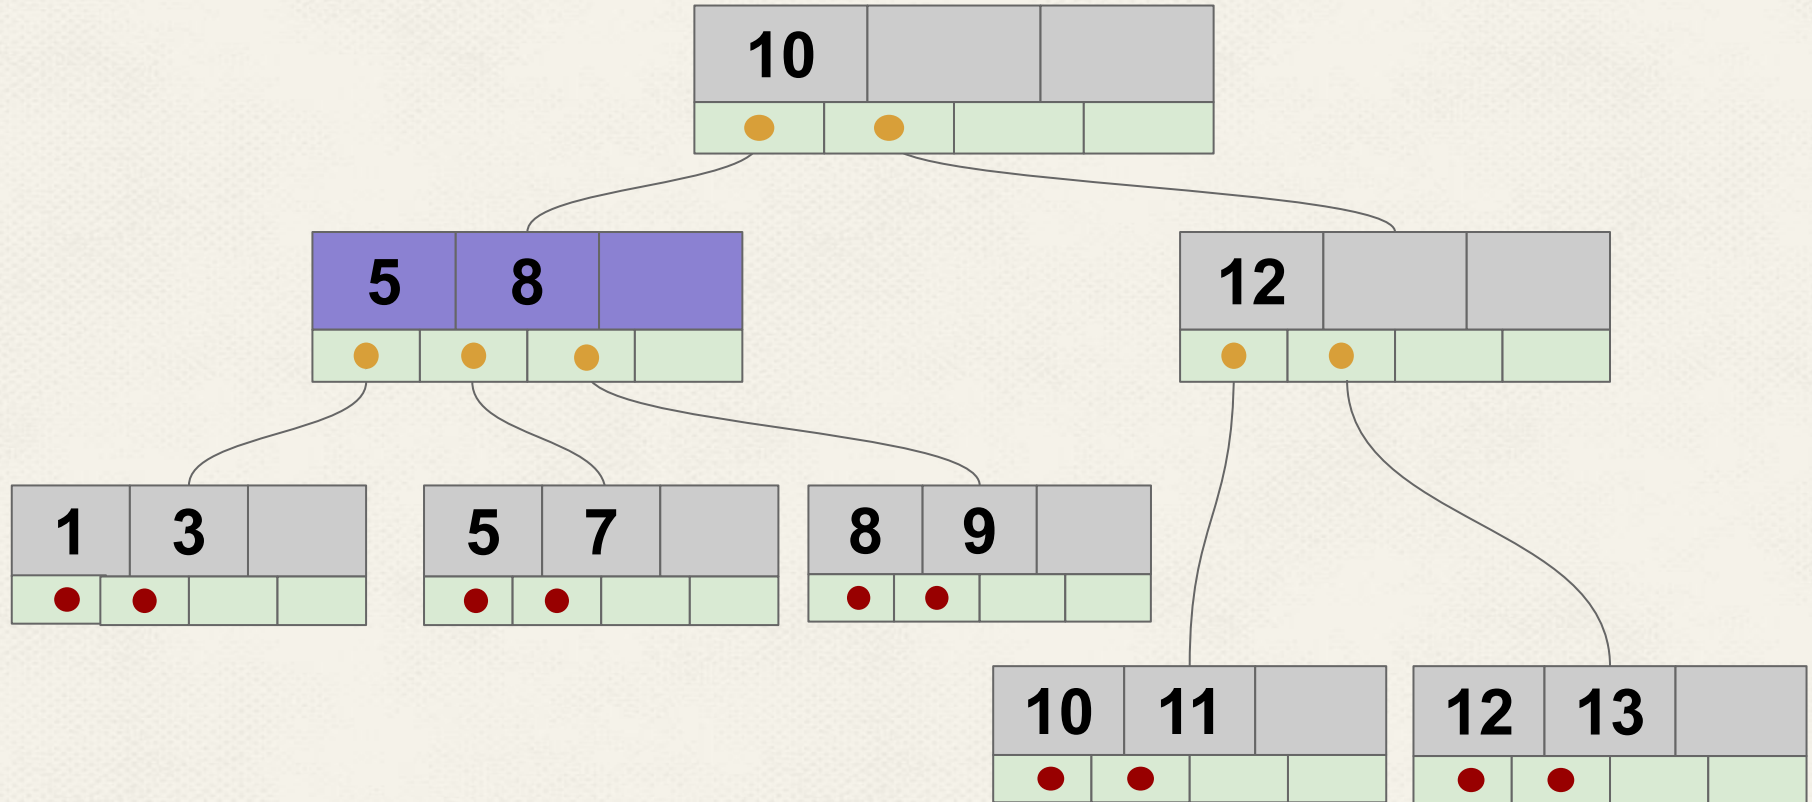
# Search Algorithm

1. Perform binary search
  - a. If the key is found on a leaf node, the data can be directly retrieved.
  - b. If the key is in an internal node, the branch must be traversed further.

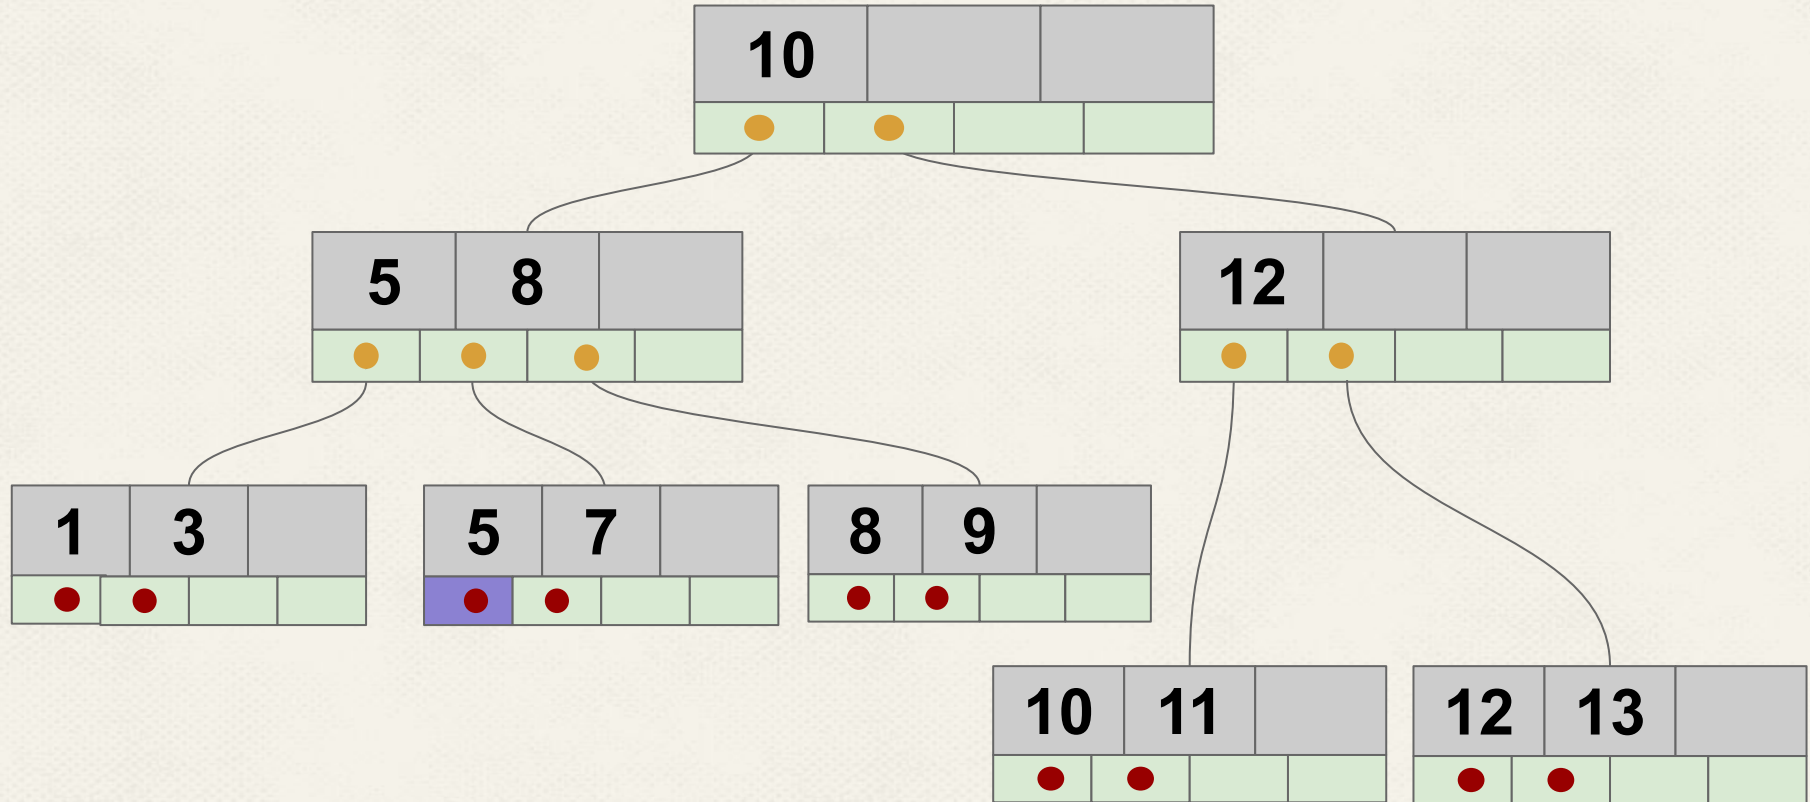
# Search for 5



# Search for 5



# Search for 5



---

# DELETE

*How do we delete a value*

# Delete Algorithm

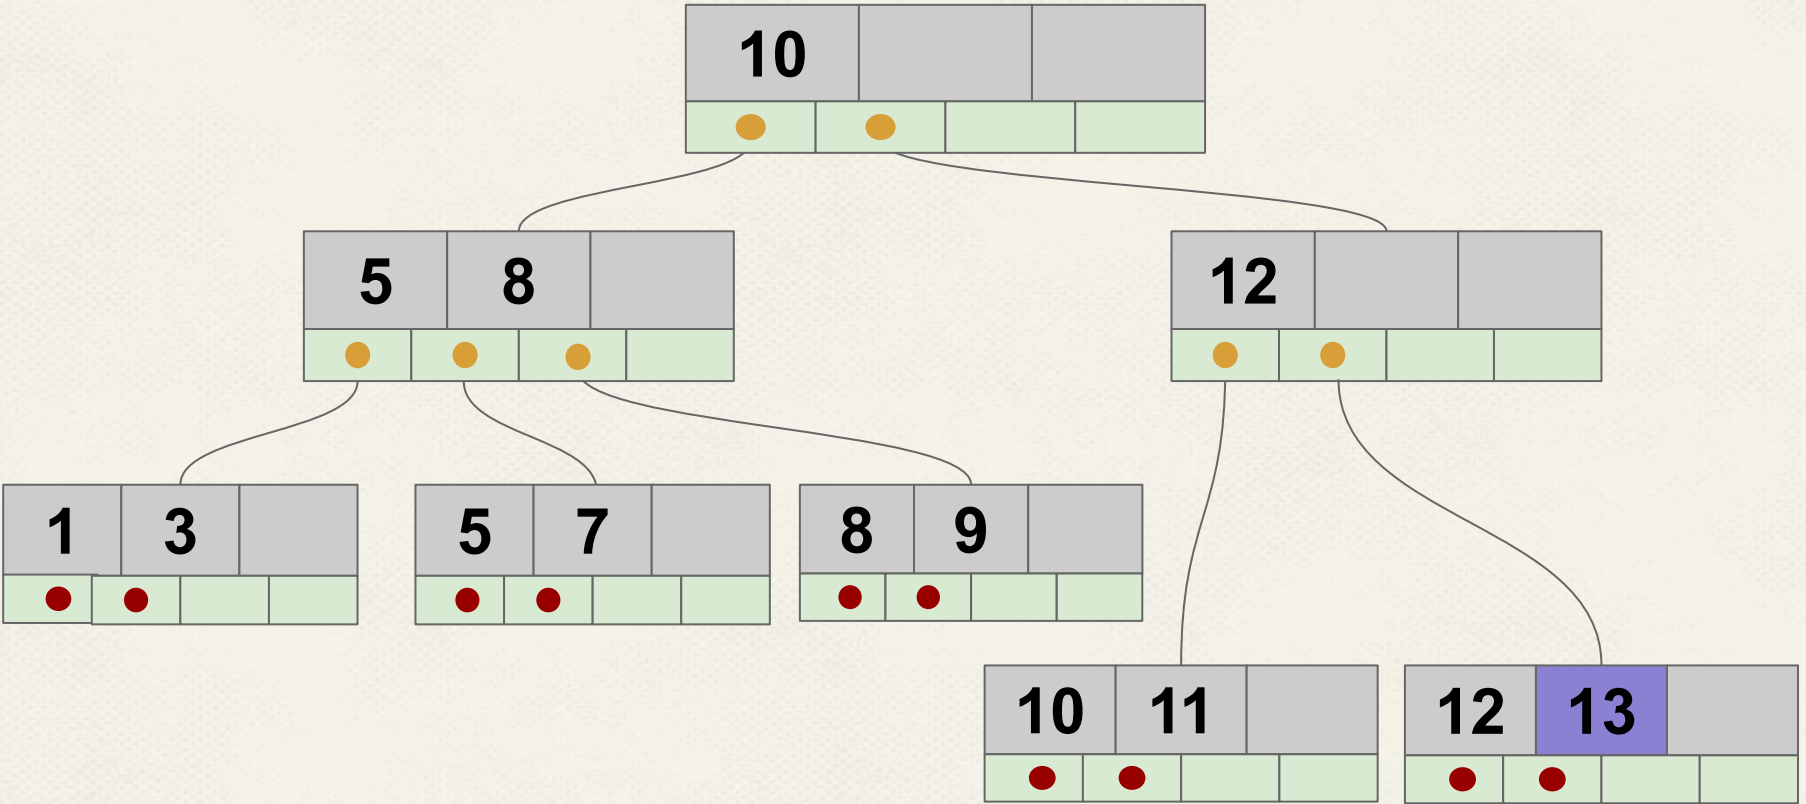
If deleting a key does not violate any of the B+ Tree properties

**Key and associated data is removed** from the node, values in the node are shifted to remove any null spaces.



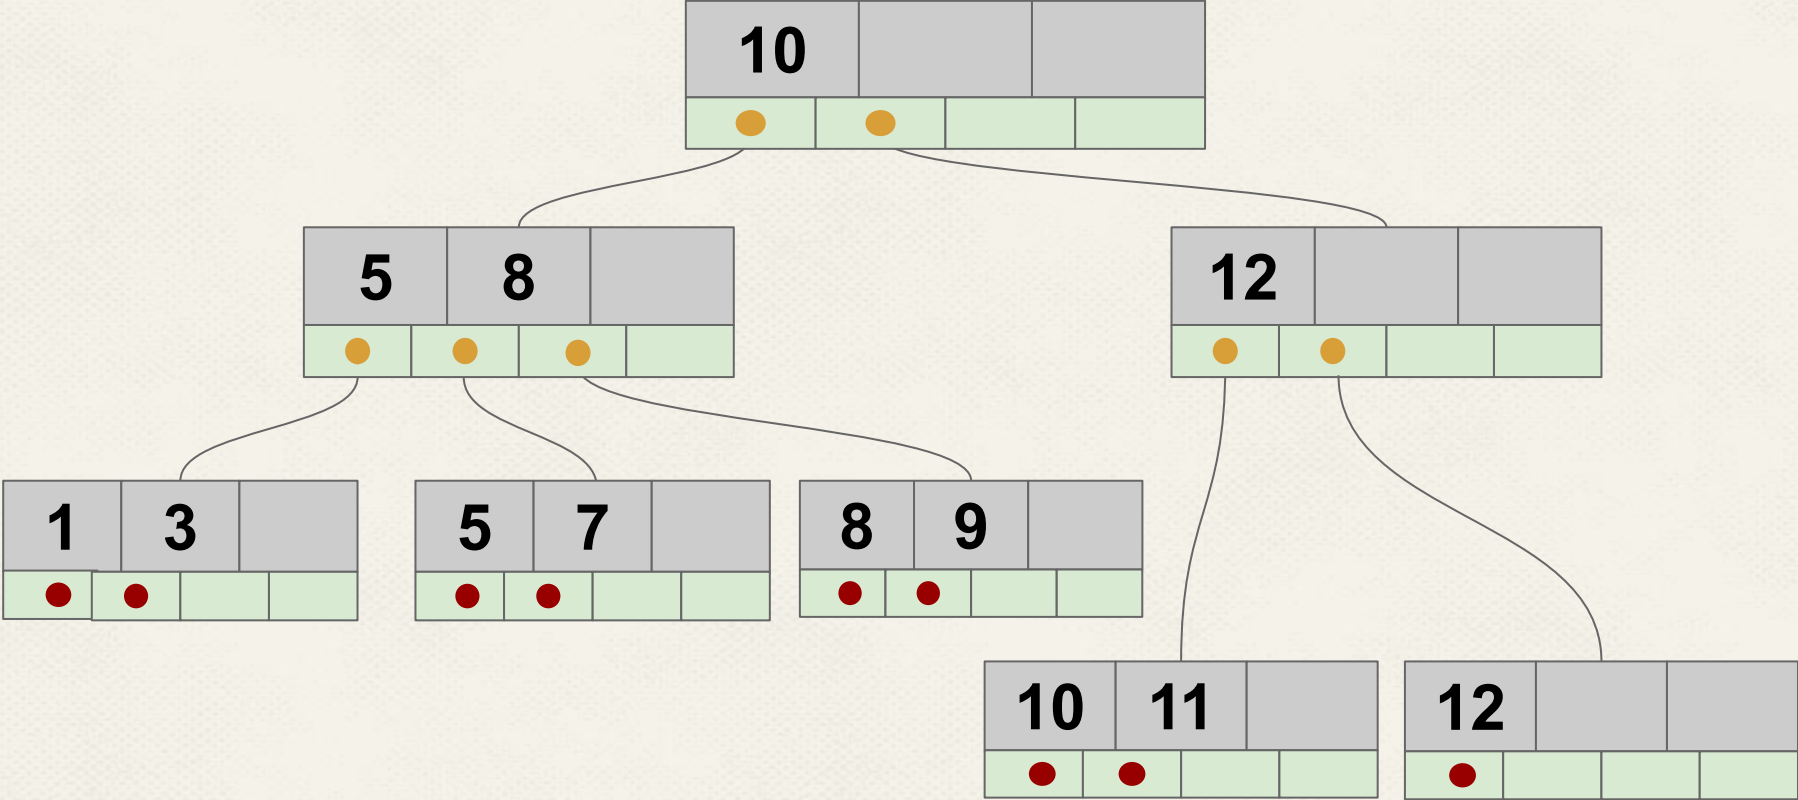
# Delete Algorithm

## Case-1: Delete 13



# Delete Algorithm

## Case-1: Delete 13



# Delete Algorithm

If the deleted key appears in both an internal node and a leaf node

The key must be deleted in the internal node must also be deleted

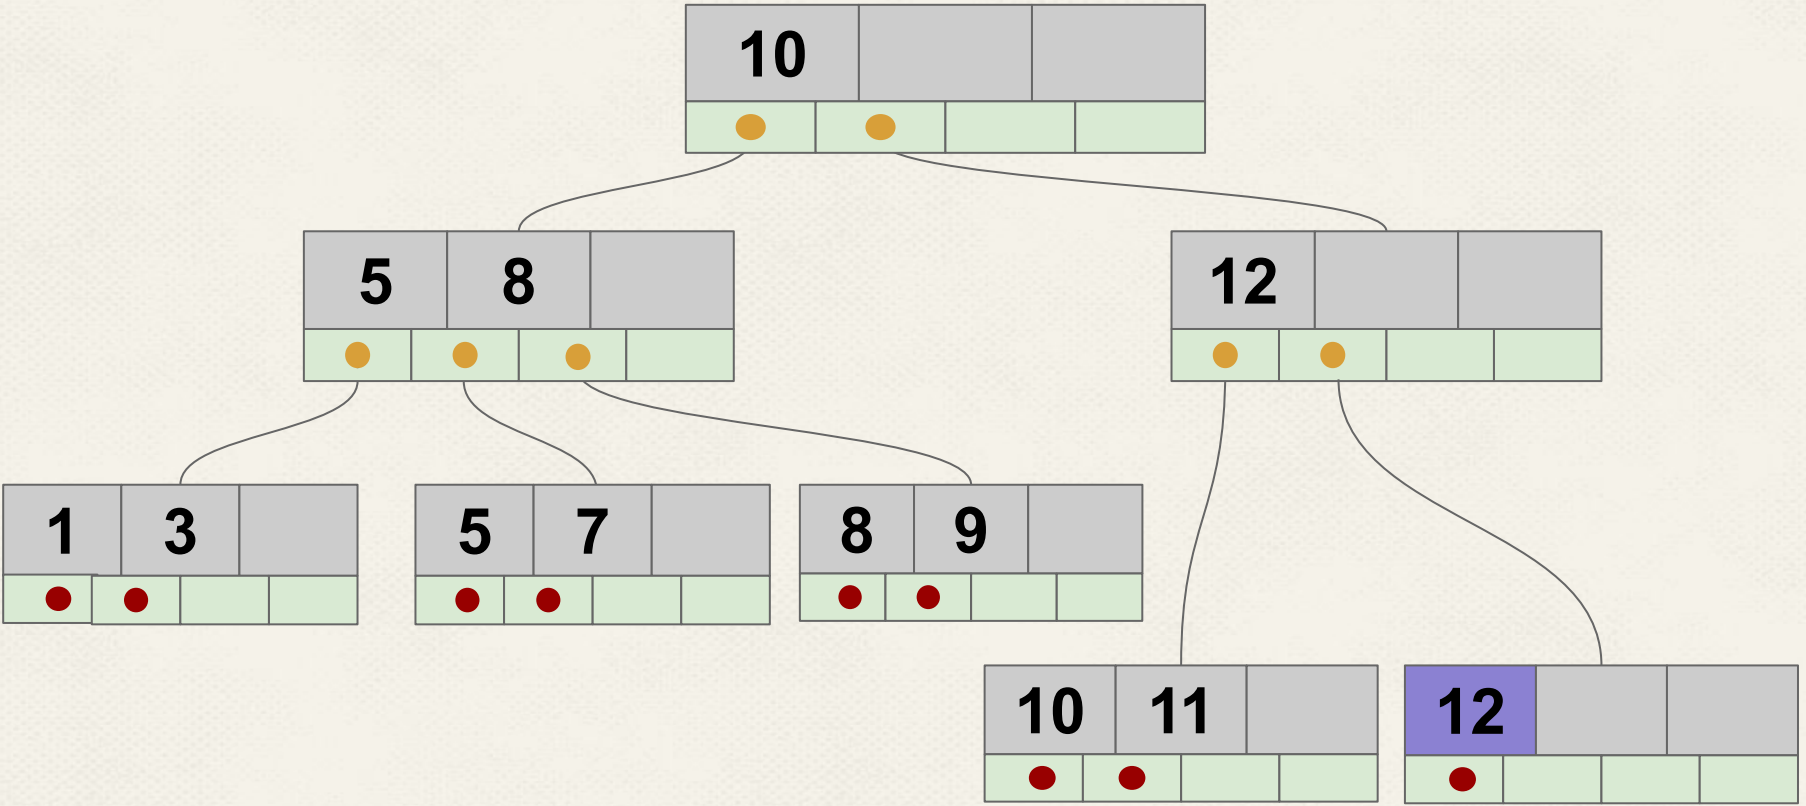
# Delete Algorithm

If deleting the key will lead to the node having fewer than  $d/2$  keys

The nodes must be **redistributed** or **merged** such that no node (except the root) ---


# Delete Algorithm

## Case-2: Delete 12




# Delete Algorithm

## Case-2: Delete 12



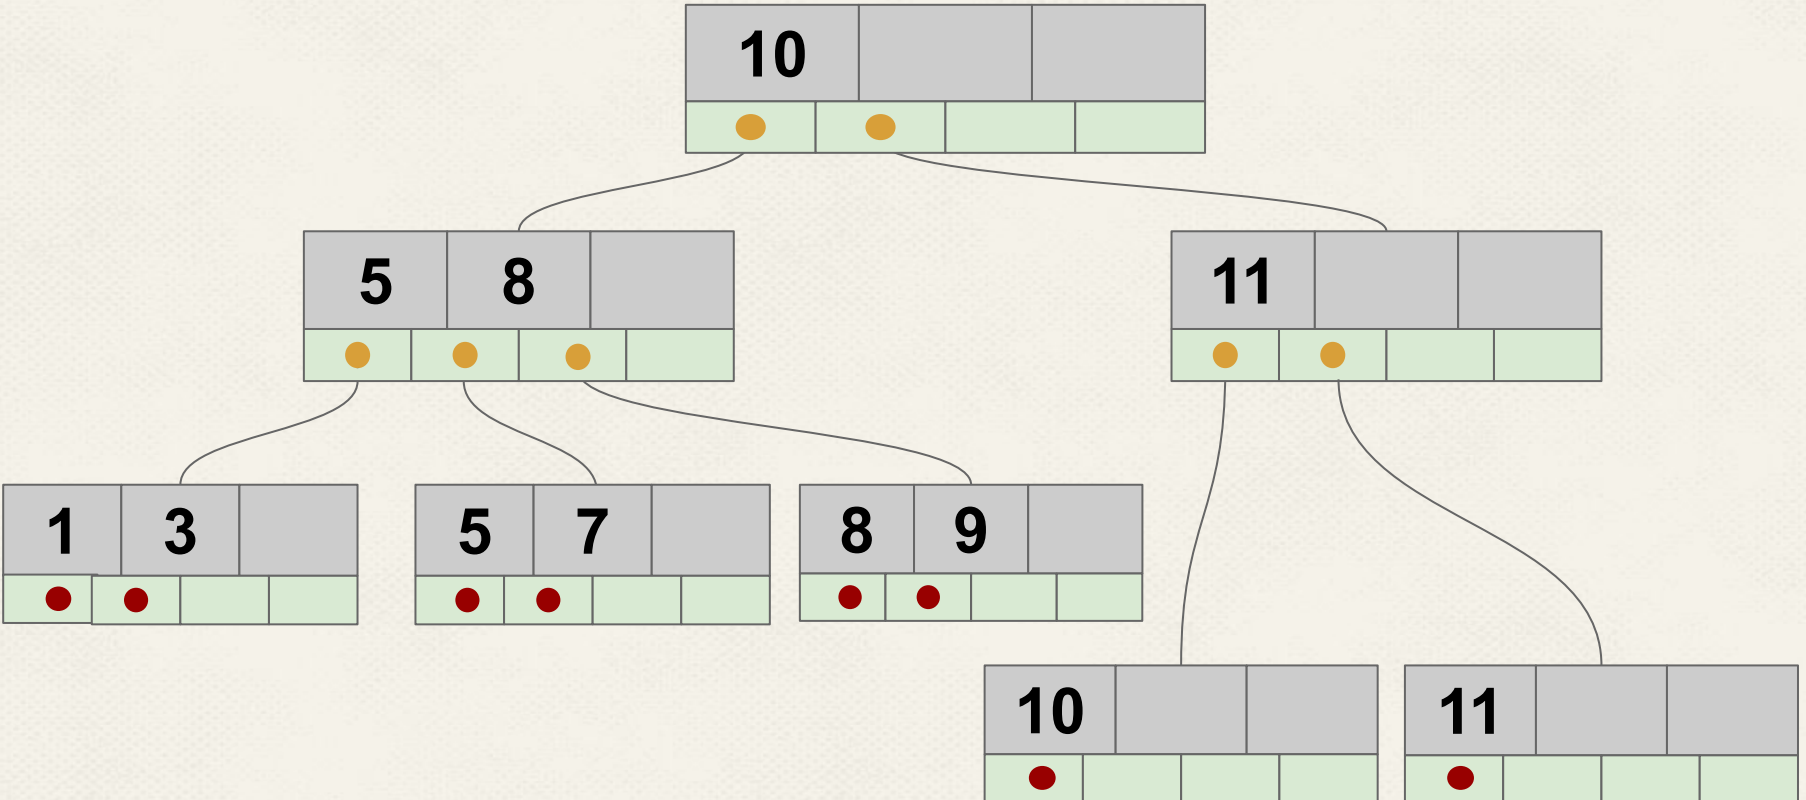
# Delete Algorithm

## Case-2: Delete 12



# Delete Algorithm


## Case-2: Delete 12





# Delete Algorithm

## Case-2: Delete 12




# **Thanks for listening!**


Any questions?




# Deletion Algorithm




## Deletion Algorithm (cont)



# Insertion Algorithm



## Insertion Algorithm (cont)



# What is a B tree?

**Self balancing binary search tree in which each node can have multiple children**

1. The keys in a B-tree are always sorted
2. Every leaf of a B-tree is on the same level
3. A child between key-1 and key-2 will have keys greater than key-1 and less than key-2



# What is a B tree?

**Self balancing binary search tree in which each node can have multiple children**

---

1. The keys in a B-tree are always sorted
2. Every leaf of a B-tree is on the same level
3. A child between key-1 and key-2 will have keys greater than key-1 and less than key-2

# What is a B tree?

**Self balancing binary search tree in which each node can have multiple children**

---

1. The keys in a B-tree are always sorted
2. Every leaf of a B-tree is on the same level
3. A child between key-1 and key-2 will have keys greater than key-1 and less than key-2

# What is a B+ Tree?

## An extension of a B-tree

1. Data can only be stored in leaf nodes.
2. Leaf nodes are sometimes linked together.

# B+ vs B Trees

1. Elements always deleted from leaf node.  
(Faster and simpler deletion)
2. Leaf nodes are linked.  
(Faster search operations)
3. Smaller height and remains balanced.