

# On Disk B+ Trees

Databases Spring 2019

# B+ Trees On Disk

- 1MB block pull from hard drive = 20ms
- B+ Tree
  - Root = 5000 indices
  - Internal Height 1 =  $5000 * 5001 = 25,005,000$  indices
  - Max 3 block pulls for 25,000,000 indices
  - Max 6 block pulls for lookup
- Binary Search
  - $\text{Log}_2 25,000,000 = 24$
  - $\text{Log}_2 5,000 = 12$
  - 12 block pulls

# B+ Trees On Disk


Super fast range queries if blocks are sorted on disk

	Random Read 4KB Blocks (MB/s)	Streaming Reads
Hard Drive	2.1	151.6
Solid State Drive	129.4	252.6




# B+ Trees On Disk


Super fast range queries if blocks are sorted on disk




# B+ Tree




# B+ Tree



# B+ Tree




# B+ Tree






# Give Left - insert 3

Capacity = 3




# Give Left - insert 3

Capacity = 3




# Give Left - insert 3


Capacity = 3




Give Right - insert 3




Give Right - insert 3




Give Right - insert 3



Split - insert 3





## Split - insert 3







## Split - insert 3



## Split - insert 3




## Split - insert 3




# Edge cases

X = Number

**\_ = Empty**




# Edge cases




# Edge cases

Parent pointers



# Edge cases

Travel down twice  
after every insert



# Rotations for insertion

- Internal -> Leaf
  - left, right, split
- Internal -> Internal
  - left, right, split
- Root -> Internal
  - split



# No Padding

```
("Good Omens", 1990, 432, "0060853980")("A Distant Mirror", 1972, 677, "0345349571")XXXXXX  
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

# End Padding

```
("Good Omens", 1990, 432, "0060853980")XXXXXX  
("A Distant Mirror", 1972, 677, "0345349571")  
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

# End Padding

```
("Good Omens", 1990, 432, "0060853980")XXXXXX  
("A Distant Mirror", 1972, 677, "0345349571")  
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

```
("Good Omens", 1990, 432, "0060853980")X  
("A Distant Mirror", 1972, 677, "0345349  
571")XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
```

# Field Padding

```
("Good OmensXXXXXX", 1990, 432, "0060853980")  
("A Distant Mirror", 1972, 677, "0345349571")  
("$100000 under th", 2009, 396, "0307274939")
```

# Strings vs Bits

```
("Good OmensXXXXXX", 1990, 432, "0060853980")  
("A Distant Mirror", 1972, 677, "0345349571")  
("$100000 under th", 2009, 396, "0307274939")
```

```
101101010101100100010101011011010101011010100  
10100010100101011010101001010101001010010  
1000100101011010101101010010101110101011000
```

# Python/C Structs

```
("Good OmensXXXXXX", 1990, 432, "0060853980")  
("A Distant Mirror", 1972, 677, "0345349571")  
("$100000 under th", 2009, 396, "0307274939")
```

```
(16 char[], int, int, 10 char[])
```

```
import struct
```

```
format_string = '16s i i 10s'
```

# Future

## Things I should do

- Finalize rotations so B+ Tree works properly
- Use a file and structs to make the B+ Tree work on disk

## Cool things I could do

- Implement delete
- Binary Search Trees instead of Lists for internal nodes
- Read through an actual implementation
  - B+ Tree resources are extremely limited