

SQL Queries

Leon Lam, Rachel Yang

A large, dark blue, abstract shape that starts from the bottom left corner and extends diagonally upwards towards the right, covering the bottom half of the slide.

How is a query processed?

It goes through **pre-optimization** (parsing, binding, other checks), **optimization** (creating and comparing query plans), and finally **execution**.

Pre-optimization

```
SQL> SELECT * FORM employees;  
SELECT * FORM employees  
      *  
ERROR at line 1:  
ORA-00923: FROM keyword not found where expected
```

Parsing

- Is the query written correctly?
- Checks syntax:
 - Spelling words wrong
 - Misusing reserved words
 - Forgetting semicolons
- Outputs:
 - Parse / sequence tree with logical execution steps

Pre-optimization

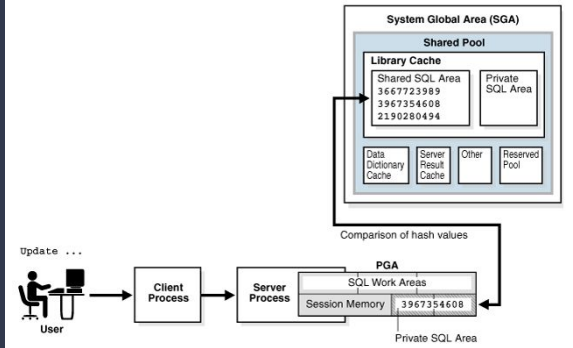
```
SQL> SELECT * FROM nonexistent_table;  
SELECT * FROM nonexistent_table  
      *  
  
ERROR at line 1:  
ORA-00942: table or view does not exist
```

Binding - performed by algebrizer

- Is the query statement meaningful — are the objects valid and referenced correctly?
- Checks semantics:
 - Existence of objects
- Outputs:
 - Query processor tree

Pre-optimization

Figure 3-2 Shared Pool Check

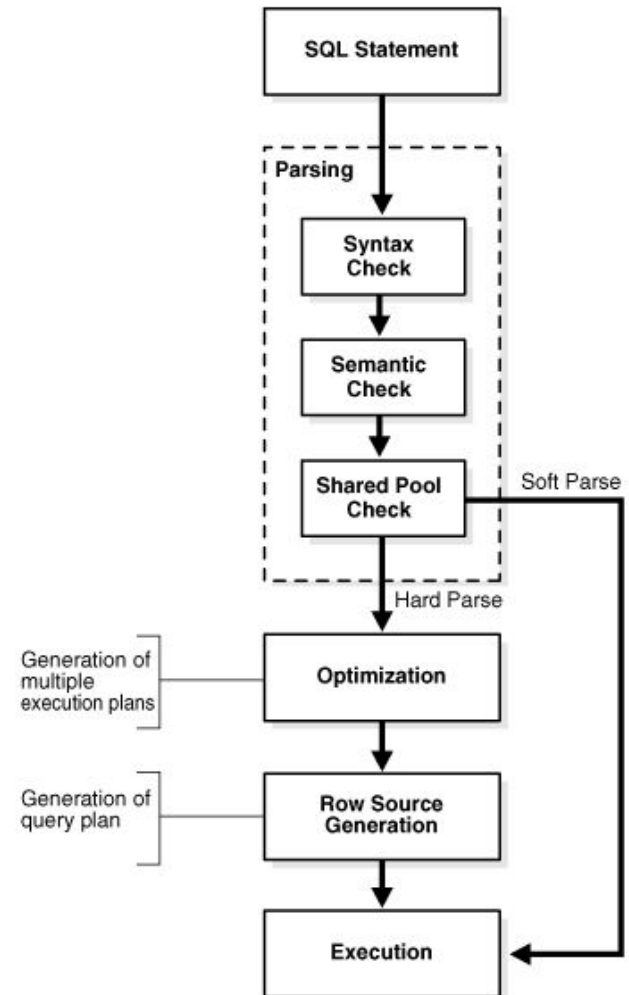


Other Checks - i.e. Shared Pool Checks

- Can the DB skip resource-intensive processing steps?
- Checks:
 - Has this query been encountered before? *By using hashing algorithm*
- Outputs:
 - Result of either soft or hard parse

Shared Pool Checks (expanded)

Figure 3-1 Stages of SQL Processing




What is query optimization?

Finding the **most efficient** way to execute a given query for a given database, by comparing possible query execution plans.

most efficient = least costly



fastest
execution time



of steps
required



and
more

How is query optimization done?

There's a few pre- and post-optimization processes, but ultimately it **looks at all (or most) possible orders** of operations, **builds trees** for each execution plan, and **compares these trees by cost**.

Optimization

Create execution plans

- How does the software know what execution plan gives us the desired result?
- Compute different permutations of operations
 - Different types of access paths
 - Different relational table join techniques / join order
 - Reconsidering indexes
- Output: Node tree where each node is a separate operation required to execute a query
 - Flows from bottom to top – child nodes output to parent nodes

Optimization

Compare execution plans

- How to compare plans across multiple cost metrics (execution time, monetary cost)?
- In each plan, each step is assigned an estimated cost value
 - Cardinality flowing through each edge in a query plan affects number of operations
 - These cardinality estimates depend on estimates of the selection factor of predicates in the query
 - * If predicates are combined, the accuracy of estimation decreases because of high correlation between predicates
- Each plans' total estimated cost is the sum of its steps' estimated costs
- Creating and comparing plans is also costly

Why is query optimization important?

Although they arrive at the same result, an inefficient query can be many times **more expensive** than an efficient one.

```

SELECT * FROM
  (SELECT * FROM INT_BIG where big_val%2 = 0) as b
  INNER JOIN
  (SELECT * FROM INT_SMALL where small_val % 2 = 0 AND small_id = 1000) as s
  ON b.big_id = s.small_id
  INNER JOIN
  (SELECT * FROM INT_MED where med_val % 2 = 0) as m
  ON m.med_id = s.small_id
;

```

```

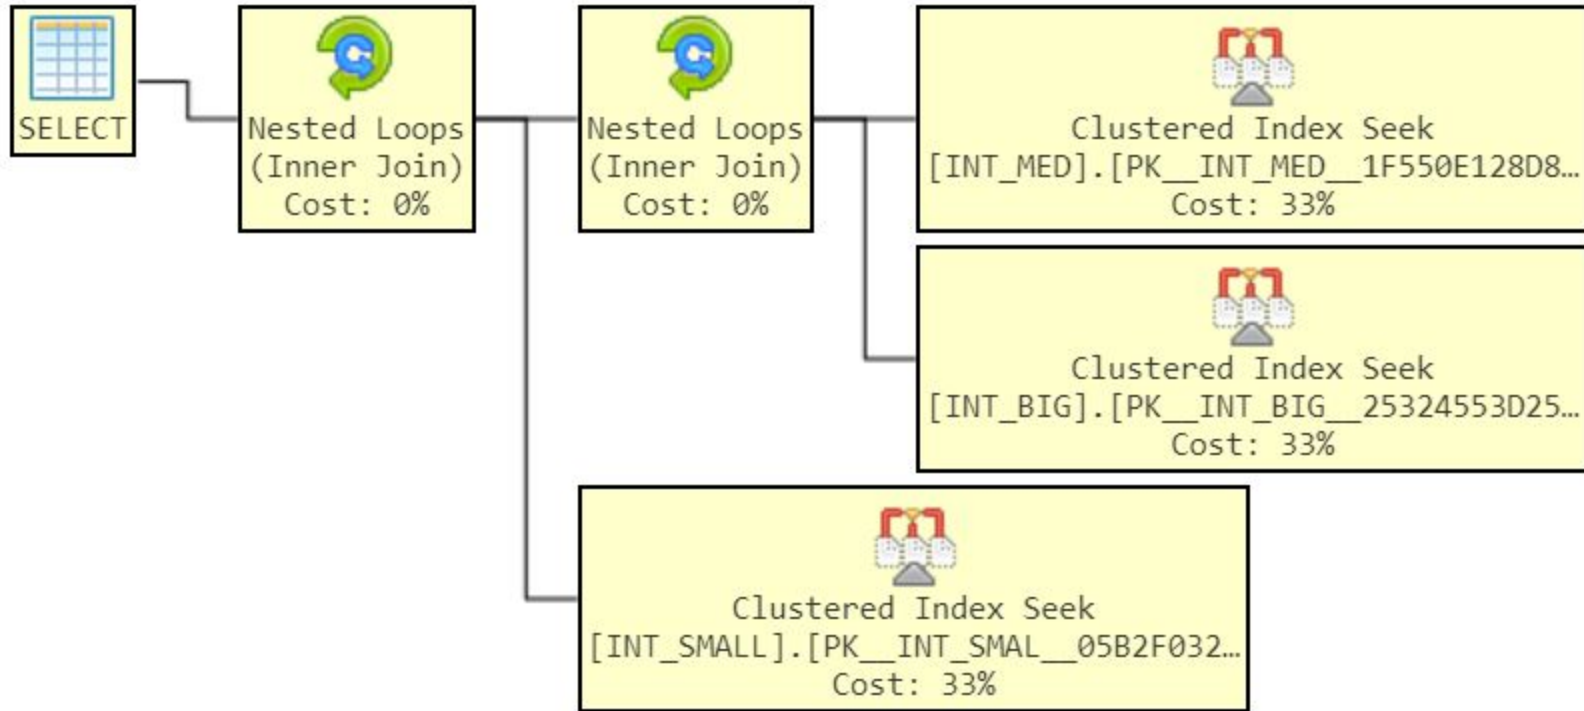
SELECT * FROM (
  SELECT * FROM INT_BIG b
  INNER JOIN INT_SMALL s ON b.big_id = s.small_id
  INNER JOIN INT_MED m ON m.med_id = s.small_id
) as res
WHERE res.big_val % 2 = 0 and res.small_val % 2 = 0 and res.med_val % 2 = 0
and res.small_id = 1000
;

```

- **Green** was written to filter the individual tables before joining
- **Red** was written to join first and then filter afterward
- If parsed sequentially, we would expect **Green** to be much faster than **Red**

Example queries

Microsoft SQL Server 2005 XML Showplan



Results: Both queries were optimized to same query plan

How does the exact optimization math work?

We have no real idea. **Each database is different** in terms of what the data looks like and which cost metric is more important.

Post-optimization

Plan Storage

- Once a plan is selected as “best,” it’s stored in a plan cache in memory
 - UNLESS an identical plan already exists in the cache
- These plans are accessed by SQL Server and then executed