

Introduction

Spring 2025

Agenda

- Why databases?
- Structure of the course
- Go

Why do we need databases?

- Applications need to manage data
 - internal data structures for storing data
 - arrays, hashmaps, dataframes, ...
- BUT: how do you persist the data structures?
 - so that they survive an application stopping and restarting
- BUT: how do you share the data structures?
 - so that more than one user can interact with the data at the same time
- BUT: how do you manage data that's larger than memory?
 - historical data
 - system logs
 - large machine learning training sets

Databases

Pull data out of the application and persist to disk

- DIY in the early days
- tends to be slow and error-prone to implement by hand

Create "generic" tool to store data on disk — a database

- Provide a convenient interface to the data
- Efficient storage and retrieval mechanism
- Run as a server accessed by other applications
- Support concurrent access and replication

Databases support a specific **data model** (how data is conceptualized)

Relational databases

Data conceptualized as rows of values in **tables**

- Think Excel sheets
- Original data model
- **Convergence** to a unified way of presenting that model in the 70s
 - Oracle, Postgres, MySQL, Microsoft SQL Server, ...
 - Standardized query language: **SQL**
- Requires pre-defined data schemas (think static typing)
- ACID properties: atomicity, consistency, isolation, durability
 - basically each user has the illusion of being the only user of the database

Non-relational databases

Also known as NoSQL databases — any database that is not relational:

- Document databases: MongoDB, CouchDB, Firebase Realtime...
- Key-value stores: Redis, DynamoDB, Cassandra, HBase, ...
- Graph databases: Neo4j, ...

They relax the structural and consistency restrictions of relational databases

- No need for schema definitions (think dynamic typing)
- More resilient replication
- What's the catch?

Looks complicated...

As a casual programmer, beyond figuring out whether you'll be storing arrays (relational) or storing JSON objects (non-relational), **it mostly doesn't matter** which database you use for your projects

They're all pretty much equivalent for casual use

There are big differences, but they don't really appear until you have millions of data points, high traffic, or extensive database replication to reduce latency

- Most likely, an architect at your company will have chosen for you

... and sometimes it is complicated

A project/startup starting small, growing huge fast

- The database choices made early often won't scale
- Need for database migration
- Made easier by abstracting the interface to the DB as much as possible
 - DAL = data access layer, which interacts with the DB directly
 - interact with the DAL via IPC / REST calls
 - migrating databases requires changing the DAL, not the client
- There are other approaches (ORMs, etc)

Structure of the course

We're going to study the various classes of databases that exist

- relational vs non-relational databases
- learn how to interact with them
- learn the key features of each class
 - data structuring
 - query languages
 - transactions, concurrency
 - replication, consistency
- learn a little bit how they work inside
 - this is not a database theory course
 - learn enough to understand what is happening
 - (learn enough to know you don't want to roll-your-own except for small projects)

Structure of the course

- Lecture-based — I like talking
 - I also like to demo code
 - Textbook *Principles of Database Management*
- Homeworks
 - Go code to build up an understanding of the concepts we study
 - some hands-on experience with databases: install them, populate them, query them
 - work in pairs (choose your team)
- Final project
 - choose project mid-semester
 - can be a deeper exploration of a topic we talked about, or an introduction to a topic we skipped
 - final presentations during events week
- Meet on Mondays, office hours before class
 - No CA, but I'm available for questions, and we can find tutors if you need one
 - Speak up, reach out — as usual, the more interactive it is the more fun it is

Introduction to Go

(online demo)