


NoSQL Databases

Spring 2025

Last time...

Distributed databases:

- fault tolerance serving data even if database instances fail
- localization keeping relevant data geographically close to users
- scalability serve more data to more users



CAP Theorem

A distributed system can only guarantee **at most 2** of the following properties at the same time:

C onsistency	all nodes have the same data
A vailability	every node can respond to any request
P artition tolerance	works even is messages are slow or dropped

P is pretty much required, because internet

So you can either guarantee C or A — but not both

Relational databases focus on CP — what about AP?

Today...

Introduction to NoSQL databases

- focus on availability over consistency
- because of that, much simpler structurally

History

Early 2000s — Web 2.0

- move from informational sites to social sites
- user generated content, participatory culture
- posts, likes, shares
- important to get right, but not *critical*
- database can be inconsistent for a period of time
- availability and high-throughput is more important
- easy scalability — add a new instance quickly if load spikes

BASE properties

ACID properties for relational database enforce transactions and consistency

NoSQL databases focus instead of the BASE properties

B asically A vailable	nodes are available, but not necessarily all data in them
S oft State	data read from any node may not be the latest
E ventual consistency	all nodes eventually hold the same data

No joins, no transactions

- any complexity is pushed to the apps using the database

Less database than data store — persistent distributed data structure

- one reason why devs like them so much?

Key-Value Stores

Dynamo, Redis

Distributed dictionary / hash map

$f : \text{key} \rightarrow \text{value}$

Operations:

- get value associated with a key
- put (associate) a value with a key

Goes all the way back to Thompson's [dbm](#) on Unix!

Key-Value Stores

Often distributed by [sharding](#)

- recall that sharding is a horizontal fragmentation mechanism

Values stored in different instances (nodes) of a KV store based on key

Naive hashing to find which node to put the value into

- $h_{\text{naive}} : \text{key} \rightarrow \{1, 2, \dots, K\}$

Adding or removing a node requires creating a new hash function

- scan and re-hash all objects!

Consistent Hashing

A clever way to hash that supports adding / removing nodes


- without needing to rehash *everything*
- used for peer-to-peer networks, content delivery networks (Akamai)
- also known as a distributed hash tables

Intuition:

Do not hash into the discrete set $\{1, \dots, k\}$

Hash into the **interval** $[0, 1]$!

Consistent Hashing




Use two hash functions:

Hash function h_1 for nodes
 $h_1(N) \in [0, 1]$

Hash function h_2 for objects
 $h_2(obj) \in [0, 1]$

obj goes into node N
where $h_1(N)$ is first clockwise
from $h_2(obj)$

Consistent Hashing



Adding a new node (N_4):

Only need to rehash the items in the node **after** the new node

Removing a node:

Only need to rehash the items in the node to be removed

Wide Column Stores

BigTable, Cassandra, HBase

Basically a two-dimensional key-value store

$f: (key_1, key_2) \rightarrow \text{value}$

- each entry has a primary key (for sharding)
- each entry has a secondary key

Can think of the primary key as a "row" name, secondary key as a "column" name

- main operation: get value in row/column cell
- different rows may have different "columns" filled

Wide Column Stores

	A	B	C	D	E	F	G	H	I
1	v_1						v_{10}	v_{12}	
2	v_2	v_4					v_{11}		v_{13}
3			v_6		v_8	v_9			
4	v_3	v_5		v_7					v_{14}

$$f(1, A) = v_2$$

$$f(3, E) = v_8$$

...

Can also retrieve by row (expensive)

$$f(4) = \{A: v_3, B: v_5, D: v_7, I: v_{14}\}$$

Wide Column Stores

Column families
(can be on their own node)

	A	B	C	D	E	F	G	H	I
1	v_1						v_{10}	v_{12}	
2	v_2	v_4					v_{11}		v_{13}
3			v_6		v_8	v_9			
4	v_3	v_5		v_7					v_{14}

$f(1, A) = v_2$
 $f(3, E) = v_8$
...

Can also retrieve by row (expensive)

$f(4) = \{A: v_3, B: v_5, D: v_7, I: v_{14}\}$

Document Stores

MongoDB, CouchDB

A key-value store where values are not opaque, but instead structured (usually JSON objects)

$f : \text{key} \rightarrow \text{JSON object}$

- objects often collected together into named collections
- but no constraints on objects in a collection

Query language:

- retrieve all objects in a collection (broadcast to all nodes!)
- retrieve all objects in a collection with a given shape

That's all, folks!