

# Relational Databases

Spring 2025

## Last time

Conceptual model

Logical model

Physical model

**What is the data about?**

How do we represent the data in a specific (kind of) database?

How is the data represented in memory or on disk?

# Last time

## Entity relationship model

- a way to analyze data that you'll be trying to represent in a database
- understand the **entities** you want to represent
- understand the **relationships** between those entities you want to capture

It is a **conceptual** model

- not tied to any specific kind of database
- how do we represent data in a specific kind of database?

# Today

Conceptual model

What is the data about?

Logical model

**How do we represent the data in a specific (kind of) database?**

Physical model

How is the data represented in memory or on disk?

# Relational databases

Many databases use a logical model based on **relations**

- relational model  $\Leftrightarrow$  relational databases
- relations  $\Leftrightarrow$  tables
- tuples  $\Leftrightarrow$  tables rows (= records)

Formalized in the 1970 by Edgar Codd, developed by IBM and Oracle

- Easy model to understand
- Supports a powerful query language (SQL)
- Flexible enough for most data uses

# Relations and tables

Mathematically, a relation  $R$  is a subset of  $D_1 \times D_2 \times \dots \times D_k$

- a set of **tuples** of the form  $(d_1, d_2, \dots, d_k)$
- $D_1, D_2, \dots, D_k$  are **domains** of the relation

Example:  $< \subseteq \mathbb{N} \times \mathbb{N} = \{ (0, 1), (0, 2), (1, 2), (0, 3), (1, 3), (2, 3), \dots \}$   
where  $(3, 4) \in <$  is usually written  $3 < 4$

A relational database represent a relation via a **table**

- each **row** is a tuple of the relation
- **columns** get names and types for convenience (= schema of the table)

# Relations and tables

Each table should define a **primary key**

- a set of columns that uniquely identifies each row

Values in a row may be blank (= **null**) except for the primary key

A table row may refer to a row in another table (via that row's primary key)

- we call that a **foreign key**
- it's an indication to the database of how we can use a field
- database can enforce foreign key constraints (value exists as a key)


## ER $\Rightarrow$ relational model: entities


Suppose we have an ER diagram capturing the structure of our data

We use tables to represent entity sets in our ER diagram

table	=	entity set
table column	=	entity attribute
table row	=	entity in a set
table primary key	=	entity set primary key







## Books

<b><u>isbn</u></b>	<b>title</b>	<b>authors</b>	<b>publisher</b>	<b>year</b>
0771595565	Rebel Angels	Robertson Davies	McMillian	1981
0316296198	The Magus	John Fowles	Little Brown & Co	1965
0670312134	Fifth Business	Robertson Davies	McMillian	1970
...	...	...	...	...

## ER $\Rightarrow$ relational model: relationships

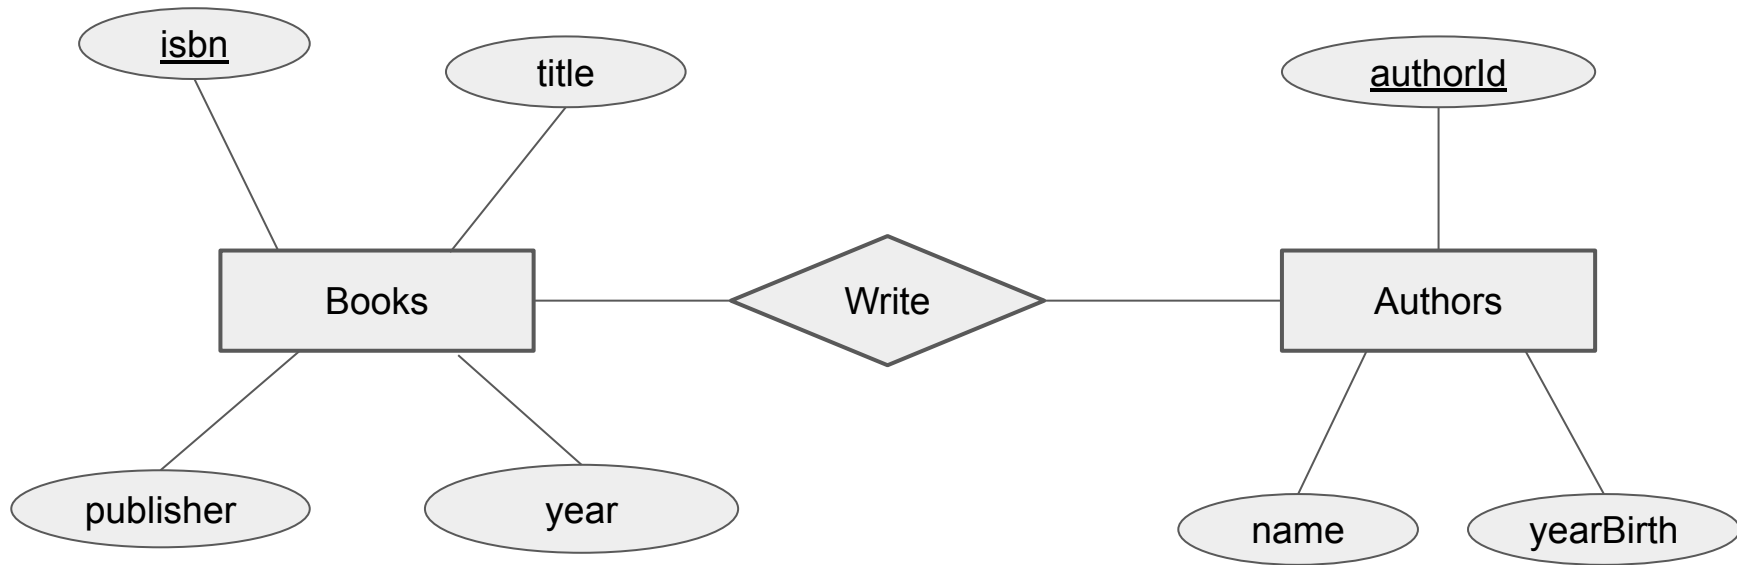
We use tables to represent relationship sets in our ER diagram

table = relationship set

table column = **foreign key** to an entity in the relationship  
also, relationship attribute

table row = relationship in a set

table primary key = entity set primary key



## Books

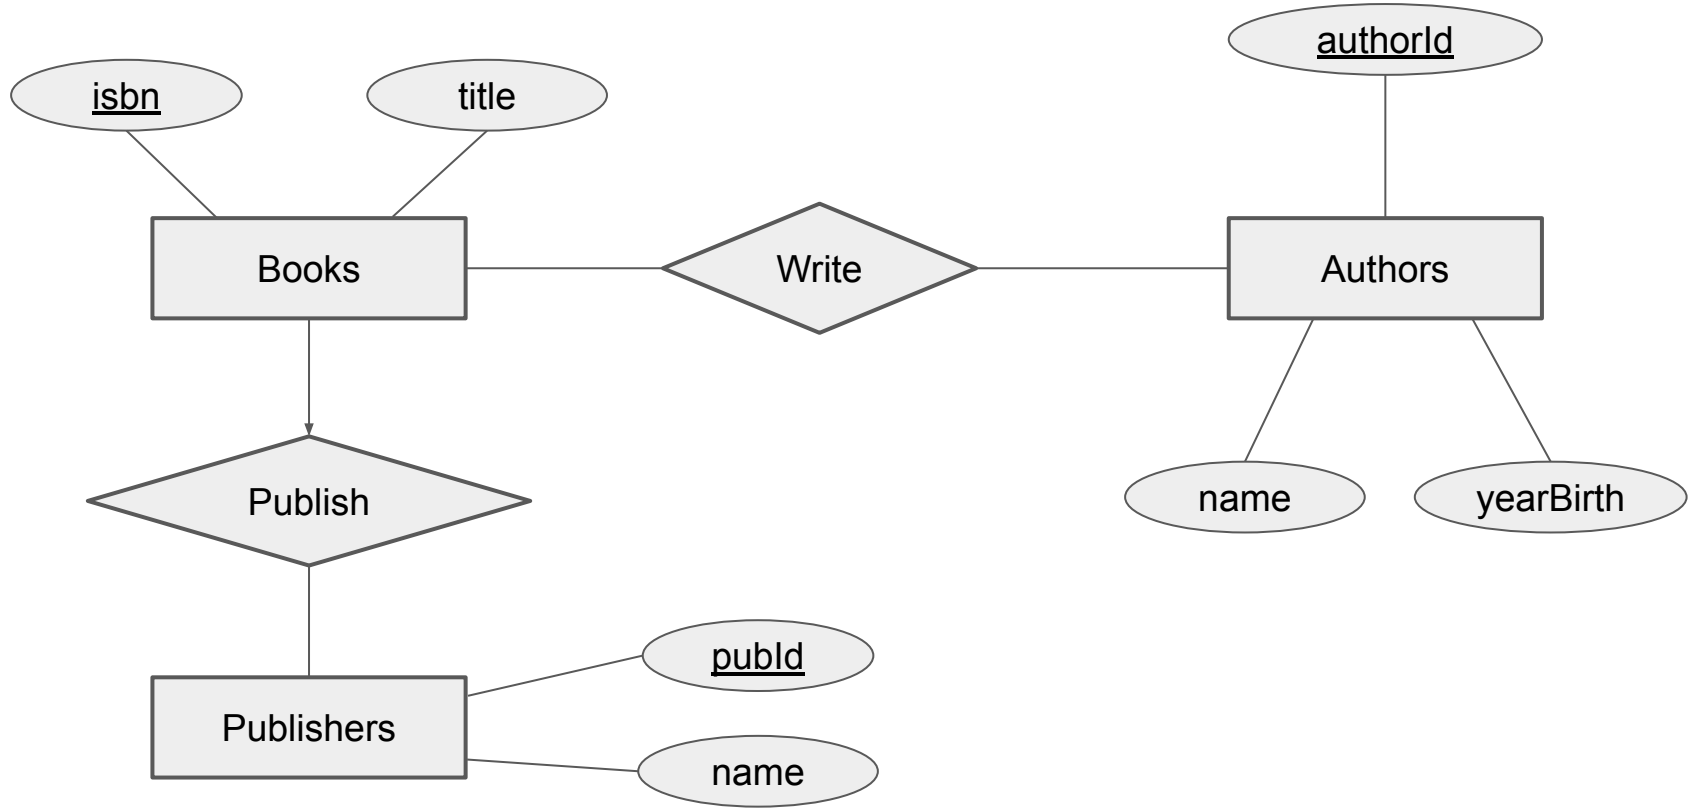
<u>isbn</u>	title	publisher	year
0771595565	Rebel Angels	McMillian	1981
0316296198	The Magus	Little Brown & Co	1965
0670312134	Fifth Business	McMillian	1970

## Authors

<u>authorId</u>	name	yearBirth
1	Robertson Davies	1913
2	John Fowles	1926

## Write

<u>isbn</u>	<u>authorId</u>
0771595565	1
0316296198	2
0670312134	1



## Books

<u>isbn</u>	title	year
0771595565	Rebel Angels	1981
0316296198	The Magus	1965
0670312134	Fifth Business	1970

## Authors

<u>authorId</u>	name	yearBirth
1	Robertson Davies	1913
2	John Fowles	1926

## Publishers

<u>pubId</u>	name
101	McMillian
102	Little Brown & Co

## Write

<u>isbn</u>	<u>authorId</u>
0771595565	1
0316296198	2
0670312134	1

## Publish

<u>isbn</u>	pubId
0771595565	101
0316296198	102
0670312134	101




## ER $\Rightarrow$ relational model: optimizing 1:N relationships

If we know a relationship is 1:1 or 1:N, we can simplify

- every book has at most one publisher

Skip the **Publish** table

Put the publisher key in the **Books** table



## Books

foreign key

<u>isbn</u>	title	year	publd
0771595565	Rebel Angels	1981	101
0316296198	The Magus	1965	102
0670312134	Fifth Business	1970	101

## Write

<u>isbn</u>	authorId
0771595565	1
0316296198	2
0670312134	1

## Authors

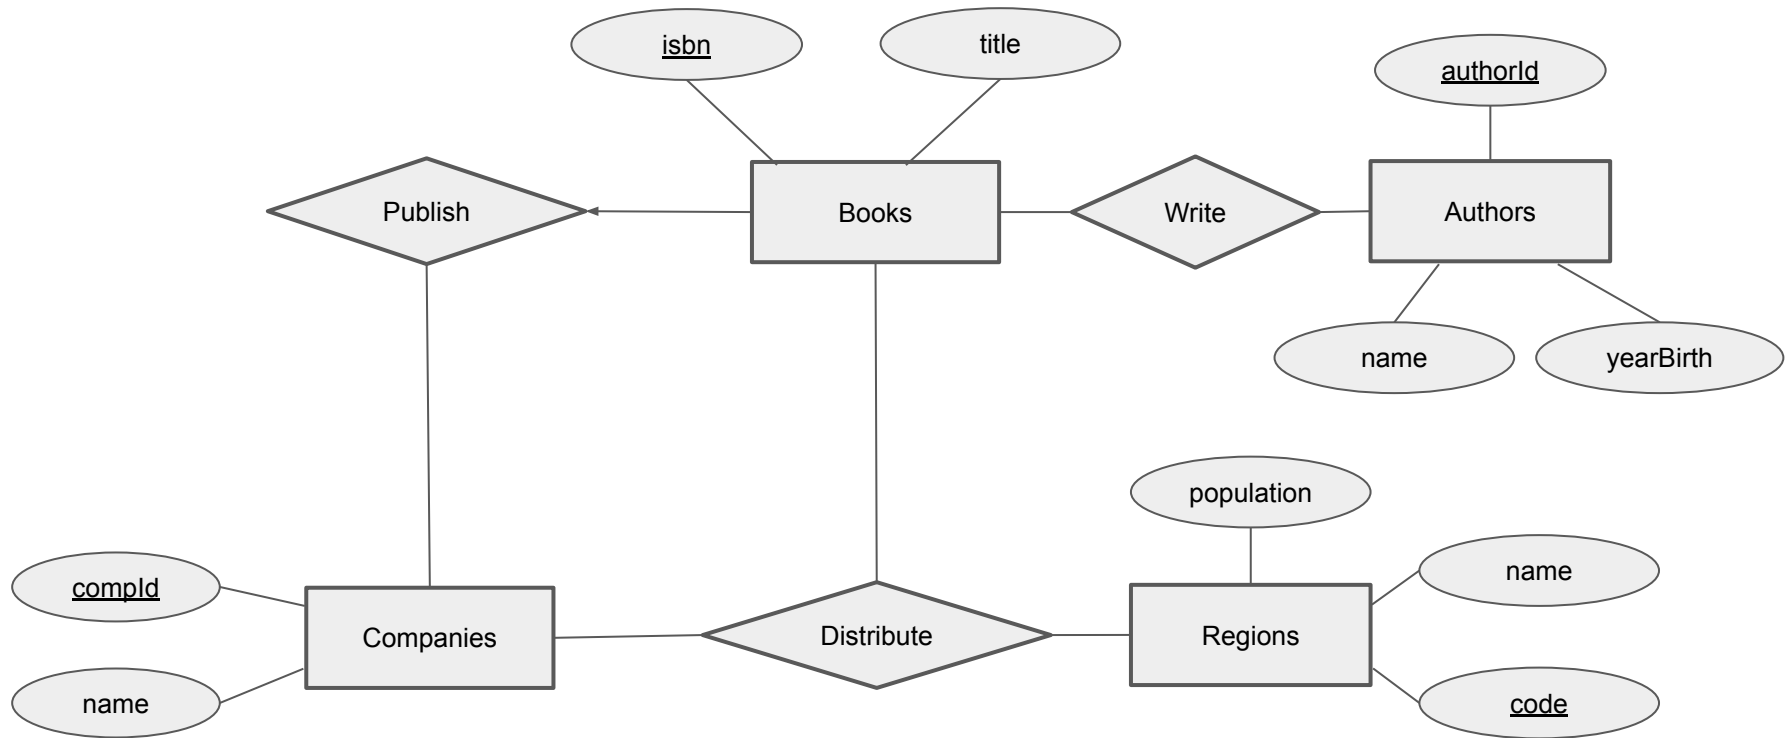
<u>authorId</u>	name	yearBirth
1	Robertson Davies	1913
2	John Fowles	1926

## Publishers

<u>publd</u>	name
101	McMillian
102	Little Brown & Co

<u>isbn</u>	publd
0771595565	101
0316296198	102
0670312134	101

Can we also remove table Write, while retaining the ability to have multiple authors per book?



Easy exercise

## In practice

How does the above work with a specific (relational) database?

- most databases run as a **server**
- use a **client** to interact with a database

We will use **SQLite**, a local *database as a library* that lives in a file on disk

- SQLite is not a server, it's a code library that uses a file as a database
- we can use SQLite interactively with the provided client
  - **sqlite3 name.db**
- a file (= a database) can host multiple tables

# In practice

Relational database API: **SQL**

- unlike most modern systems, the API is not a set of endpoints
- a string representing instructions sent to the DB

All relational databases support SQL to (1) manipulate tables/data (2) query data

- much more on querying data with SQL next time

# SQL operations on tables

**DDL** (data definition language): subset of SQL for table CRUD operations

Create a table:

```
CREATE TABLE {name} (...)
```

Update a table (structurally — add column, etc)

```
ALTER TABLE {name} ...
```

Delete a table:

```
DROP TABLE {name}
```

# SQL operations on rows

Add a row to a table:

```
INSERT INTO {table} VALUES ({value1}, {value2}, ...)
```

Update rows in a table

```
UPDATE {table} SET {field} = {value} WHERE {row condition}
```

Delete rows in a table

```
DELETE FROM {table} WHERE {row condition}
```

Read rows from a table

```
SELECT {f1}, {f2}, ... FROM {table}  
SELECT {f1}, {f2}, ... FROM {table} WHERE {row condition}
```

That's all, folks!