

SQL

Spring 2025

Last time

Conceptual model

What is the data about?

Logical model

How do we represent the data in a specific (kind of) database?


Physical model

How is the data represented in memory or on disk?

Last time

Relational databases:

- records in tables
- easy to understand
- flexible enough for most data uses
- supports a powerful query language (SQL)



Books

foreign key

<u>isbn</u>	title	year	publd
0771595565	Rebel Angels	1981	101
0316296198	The Magus	1965	102
0670312134	Fifth Business	1970	101

Write

<u>isbn</u>	authorId
0771595565	1
0316296198	2
0670312134	1

Authors

<u>authorId</u>	name	yearBirth
1	Robertson Davies	1913
2	John Fowles	1926

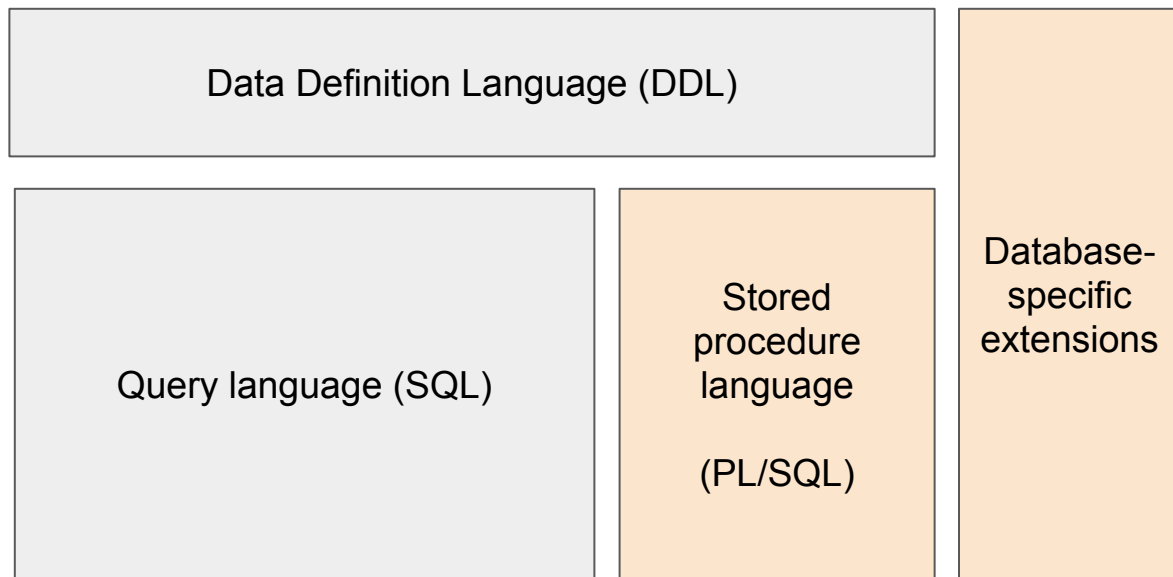
Publishers

<u>publd</u>	name
101	McMillian
102	Little Brown & Co

<u>isbn</u>	publd
0771595565	101
0316296198	102
0670312134	101

Can we also remove table Write, while retaining the ability to have multiple authors per book?

SQL at a glance



Data Definition Language

Subset of SQL for table CRUD operations

Create a table:

```
CREATE TABLE name (...)
```

Update a table (structurally — add column, etc)

```
ALTER TABLE name ...
```

Delete a table:

```
DROP TABLE name
```

Query language

Add a row to a table:

```
INSERT INTO table VALUES (value1, value2, ...)
```

Update rows in a table

```
UPDATE table SET column = value WHERE row condition
```

Delete rows in a table

```
DELETE FROM table WHERE row condition
```

Read rows from a table

```
SELECT column1, column2, ... FROM table
```

```
SELECT column1, column2, ... FROM table WHERE row condition
```


Querying via SELECT

Querying = extracting information out of the data that's stored in the database

Key operation for relational databases, and where SQL puts most of its focus

- store data on a per-table basis
- retrieve data across tables

Keep as much of the processing on the database

- you should filter / massage the data on the DB side as much as possible
- (imagine having to pull TBs of data just to access a single row)

Syntax

https://docs.aws.amazon.com/redshift/latest/dg/r_SELECT_synopsis.html



```
[ WITH with_subquery [, ...] ]  
SELECT  
[ TOP number | [ ALL | DISTINCT ]  
* | expression [ AS output_name ] [, ...] ]  
[ EXCLUDE column_list ]  
[ FROM table_reference [, ...] ]  
[ WHERE condition ]  
[ [ START WITH expression ] CONNECT BY expression ]  
[ GROUP BY ALL | expression [, ...] ]  
[ HAVING condition ]  
[ QUALIFY condition ]  
[ { UNION | ALL | INTERSECT | EXCEPT | MINUS } query ]  
[ ORDER BY expression [ ASC | DESC ] ]  
[ LIMIT { number | ALL } ]  
[ OFFSET start ]
```

Understanding queries

Four basic operations:

1. projecting
2. filtering
3. joining
4. aggregating

These operations form the basis of so-called **relational algebra** that formalizes the concept of queries in the context of relations

- relational algebra = mathematical basis of SQL queries

Intuition

A query **pulls rows from one or more tables** and **returns the rows of a new (virtual) table** representing the result of the query

- that result table is transitory and not persisted on disk
- a query just returns rows to the client issuing the query

That means that you can generally put a query wherever SQL expects a table

- nested queries

You can also create an actual table out of the results of the query

CREATE TABLE *table* **AS** (*query*)

Projecting

Return **narrower** rows from a source table

```
SELECT column1, column2, ...  
FROM table
```

Same number of rows as the source table

Special case (all columns):

```
SELECT *  
FROM table
```

Rename columns in the result

```
SELECT column1 AS name1, ... FROM table
```

Filtering

Happens before projecting

Remove some of the rows from the source table

SELECT *column*₁, ... **FROM** *table* **WHERE** *row condition*

where *row condition* is a condition that is true or false of a row of the source table

There is a full language of conditions

- *column* = *value*, *column*₁ = *column*₂, *column* > *v*
- more generally: *exp*₁ *op* *exp*₂
- *column* **IS NULL**, *column* **IS NOT NULL**
- Boolean combinations: **AND**, **OR**, **NOT**

(Fun fact)

What if you wanted to project *before* filtering?

- usually doesn't make a difference, but for the sake of argument...

Use a nested query:

```
SELECT * FROM (  
    SELECT column1, column2 FROM table  
)  
WHERE row condition
```

Joining

Until now, I've use a single source table for queries

We can also **pull rows from multiple tables** in a query

- what does that even mean?
- (conceptually) **join the tables into a single table** before querying from it
- happens before filtering and projecting

Join the tables?

- many possible interpretations

Joining — Cartesian joins

Simplest join — consider all possibilities

```
SELECT *  
FROM Books, Authors
```

- take every possible row from A and every possible row from B
- concatenate each combination into a (longer) row
- total number of rows in result = $|tableA| |table B|$

Generalizes to an arbitrary number of tables:

```
SELECT *  
FROM Books, Write, Authors
```

Example (Cartesian join)

T
A
B
C

U
X
Y

SELECT *
FROM T, U



A	X
A	Y
B	X
B	Y
C	X
C	Y

Joining — inner joins

Almost always, you want to **match** the rows from the tables

- a row from table A has a foreign key into a row of table B
- *"instead of going from table A row to the table B row, attach B row to A row"*

You can get that with filtering:

```
SELECT Books.title, Authors.name
```

```
FROM Books, Authors, Write
```

```
WHERE Books.isbn = Write.isbn AND Authors.authorId = Write.authorId
```

Same row may be matched multiple times (book with multiple authors)

Joining — inner joins

Alternative syntax:

```
SELECT Books.title, Authors.name  
FROM Books  
[INNER] JOIN Write ON Books.isbn = Write.isbn  
[INNER] JOIN Authors ON Authors.authorId = Write.authorId
```

Special case (**natural join**):

```
SELECT Books.title, Authors.name  
FROM Books  
[INNER] JOIN Write USING (isbn)  
[INNER] JOIN Authors USING (authorId)
```

Example (inner join)

T	j
A	1
B	2
C	

k	U
1	X
1	Y
	Z

SELECT *
FROM T
JOIN U on T.j = U.k



A	1	1	X
A	1	1	Y

Joining — outer joins

if a **JOIN** field value does not match for a row, the row will not show up in the **JOIN**

- this includes null values (not equal to any other value)

What if you wanted unmatched rows to be included?

- what do you join it with? nothing → other columns are nulls
- which **unmatched rows do you want to keep?**
- depends on the order in which you join

LEFT [OUTER] JOIN	→	keep unmatched rows in left table
RIGHT [OUTER] JOIN	→	keep unmatched rows in right table
FULL [OUTER] JOIN	→	keep unmatched rows in both left and right tables

Example (left outer join)

T j	
A	1
B	2
C	

k U	
1	X
1	Y
	Z

SELECT *
FROM T
LEFT JOIN U on T.j = U.k



A	1	1	X
A	1	1	Y
B	2		
C			

Example (right outer join)

T	j
A	1
B	2
C	

k	U
1	X
1	Y
	Z

SELECT *
FROM T
RIGHT JOIN U on T.j = U.k



A	1	1	X
A	1	1	Y
			Z

Example (full outer join)

T	j
A	1
B	2
C	

k	U
1	X
1	Y
	Z

SELECT *
FROM T
FULL JOIN U on T.j = U.k



A	1	1	X
A	1	1	Y
B	2		
C			
			Z

Aggregating

Done last, after projection

Aggregation → **summarize multiple rows into a single row**

- apply an aggregation function for each column

Aggregation functions:

COUNT
MIN
MAX

SUM
AVG

LISTAGG

SELECT COUNT(title), AVG(pages) FROM Books

→ returns a single row

Aggregating by groups

Can **group** rows together by one or more fields and aggregate **within** each group

```
SELECT COUNT(title), AVG(pages)  
FROM Books  
GROUP BY year
```

Can put the grouping fields in the SELECT line **without** aggregating them:

```
SELECT COUNT(title), AVG(pages), year  
FROM Books  
GROUP BY year
```

Aggregating works well with joining


```
SELECT Authors.name, COUNT(Books.title), AVG(Books.pages)
FROM Books
JOIN Write USING (isbn)
JOIN Authors USING (authorId)
WHERE Books.year > 1970
GROUP BY Authors.name
```

Can further **filter groups before aggregating** them:

...

```
HAVING COUNT(Books.title) > 2
```

Query flow



That's all, folks!

Appendix

Views

— **CREATE VIEW** *name* **AS** *query*

Distinct as a kind of filtering

— **SELECT DISTINCT** year **FROM** Books

Ordering

— **SELECT * FROM** Books **ORDER BY** year, title

Pagination: offset, limit

— **SELECT * FROM** Books **LIMIT** 10

— **SELECT * FROM** Books **OFFSET** 10 **LIMIT** 10

Table name abbreviations

— **SELECT** B.title **FROM** Books B