

Balanced Binary Search Trees

DSA, Fall 2022

Recall — Dynamic Set ADT

```
type Set  
type Cell
```

```
NewSet:    ()          → *Set  
Search:    (*Set, int) → *Cell  
Insert:    (*Set, *Cell) → ()  
Delete:    (*Set, *Cell) → ()  
Minimum:   *Set        → *Cell  
Maximum:   *Set        → *Cell
```

Recall — binary search trees

Binary trees where every node has the binary search property:

- all nodes in the left subtree of a node have value less than that node
- all nodes in the right subtree of a node have value greater than that node

To search:

- start at the root, and navigate left or right at every node based on value

To insert:

- start at the root, and navigate left or right at every node based on value
- add new node as a leaf where search would find it

Asymptotic running times

	Linked list	Doubly-linked list	Sorted doubly-linked list (with tail pointer)	Binary search tree
Search	$\Theta(n)$	$\Theta(n)$	$\Theta(n)$	$\Theta(\text{height}(n))$
Insert	$\Theta(1)$	$\Theta(1)$	$\Theta(n)$	$\Theta(\text{height}(n))$
Delete	$\Theta(n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(\text{height}(n))$
Minimum	$\Theta(n)$	$\Theta(n)$	$\Theta(1)$	$\Theta(\text{height}(n))$
Maximum	$\Theta(n)$	$\Theta(n)$	$\Theta(1)$	$\Theta(\text{height}(n))$

Asymptotic running times

	Linked list	Doubly-linked list	Sorted doubly-linked list (with tail pointer)	Binary search tree
Search	$\Theta(n)$	$\Theta(n)$	$\Theta(n)$	$\Theta(n)$
Insert	$\Theta(1)$	$\Theta(1)$	$\Theta(n)$	$\Theta(n)$
Delete	$\Theta(n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(n)$
Minimum	$\Theta(n)$	$\Theta(n)$	$\Theta(1)$	$\Theta(n)$
Maximum	$\Theta(n)$	$\Theta(n)$	$\Theta(1)$	$\Theta(n)$

Balanced binary search trees

The **height of a tree** is the number of nodes on the longest path from the root

An empty tree has height 0

A binary tree is **balanced** if its height is $\Theta(\log_2 n)$

A balanced binary search tree by definition has $\Theta(\log_2 n)$ operations

How do we ensure binary search trees are balanced?

- insert / delete as usual
- repair the tree to restore balance

Balanced binary search trees

The **height of a tree** is the number of nodes on the longest path from the root


An empty tree has height 0

A binary tree is **balanced** if its height is $\Theta(\log_2 n)$

A balanced binary search tree by definition has height $\Theta(\log_2 n)$

How do we ensure binary search trees are balanced?

- insert / delete as usual
- repair the tree to restore balance



This definition is not
really actionable

Balanced binary search trees

Many types of balanced binary search trees have been defined:

- AVL trees

- 2-3 trees

- Red-black trees

- Splay trees

- ...

Each enforces a specific property that implies that a tree is balanced

AVL trees

An AVL tree is a binary search where every node has the **AVL property**:

The difference between the height of the left subtree of a node and the height of the right subtree of a node is at most 1

This implies that AVL trees are balanced


Theorem: An AVL tree of height H has at least $2^{H/2-1}$ nodes

Corollary: The height of an AVL tree with N nodes is at most $2 + 2 \log_2 N$

Thus, the height of an AVL tree is $\Theta(\log_2 n)$

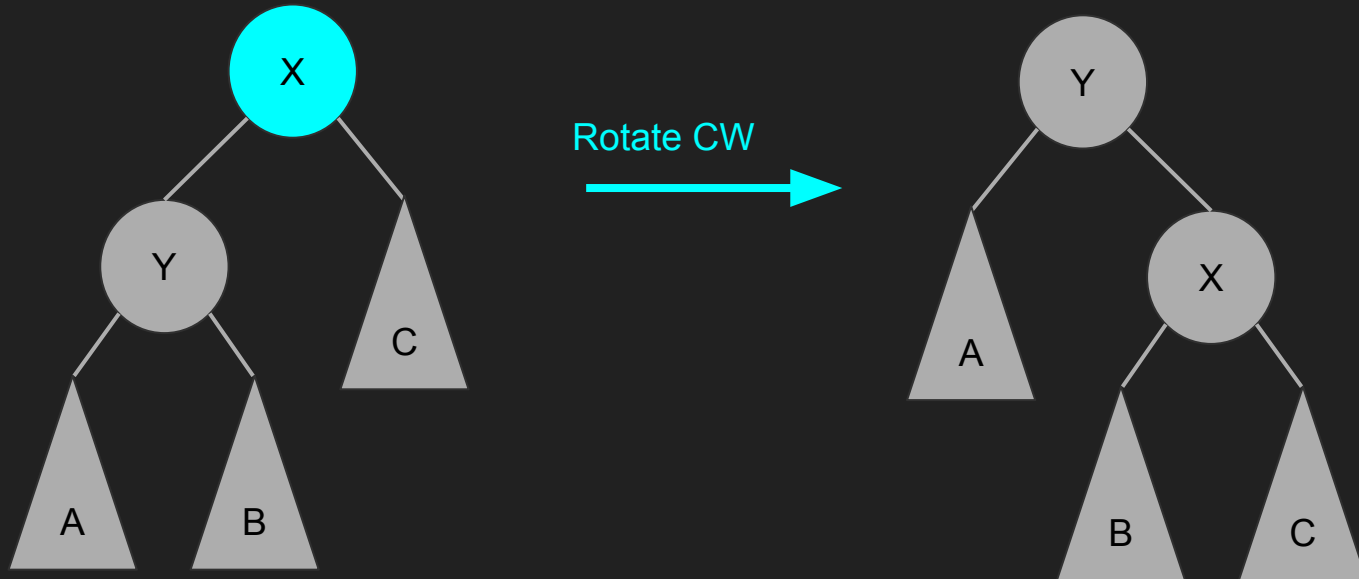
Binary search tree rotations

Rotations that preserve the BST property of all nodes



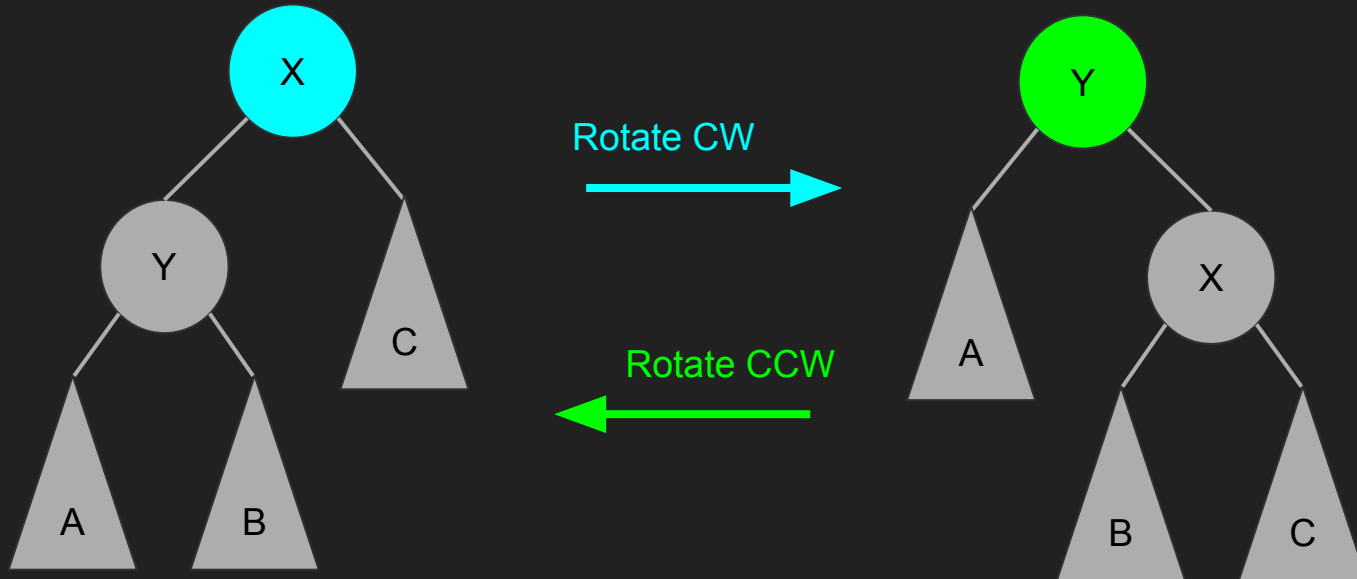
AVL trees — repair transformations

Rotations that preserve the BST property of all nodes



AVL trees — repair transformations

Rotations that preserve the BST property of all nodes



AVL trees — Insert operation

Insert V using the usual binary search tree insertion:

- start at the root
- if V is at the node you're at, you're inserting an existing value — STOP
- if V is less than the value at the node, go to the left child
- if V is more than the value at the node, go to the right child
- repeat until there are no more nodes to follow
- insert a node as the left or right child (depending on V) of the last node visited

Then:

- start at inserted node
- walk back up to the root, repairing the tree on the way up

AVL trees — repairing node X after Insert

Check height of X left and right subtrees

- if difference in height at most 1, continue up
- if difference is 2, repair then continue up

Two cases:


- Case L: highest subtree is the left subtree
- Case R: highest subtree is the right subtree

AVL trees — repairing node X after Insert (case L)


Height of left subtree of X is $h + 2$, height of right subtree of X is h

So left child of X is a node Y whose height of its subtrees are h and $h + 1$

Subcase LL:




Subcase LR:




height h

AVL trees — repairing node X after Insert (subcase LL)




height h

AVL trees — repairing node X after Insert (subcase LL)




height h

AVL trees — repairing node X after Insert (subcase LR)



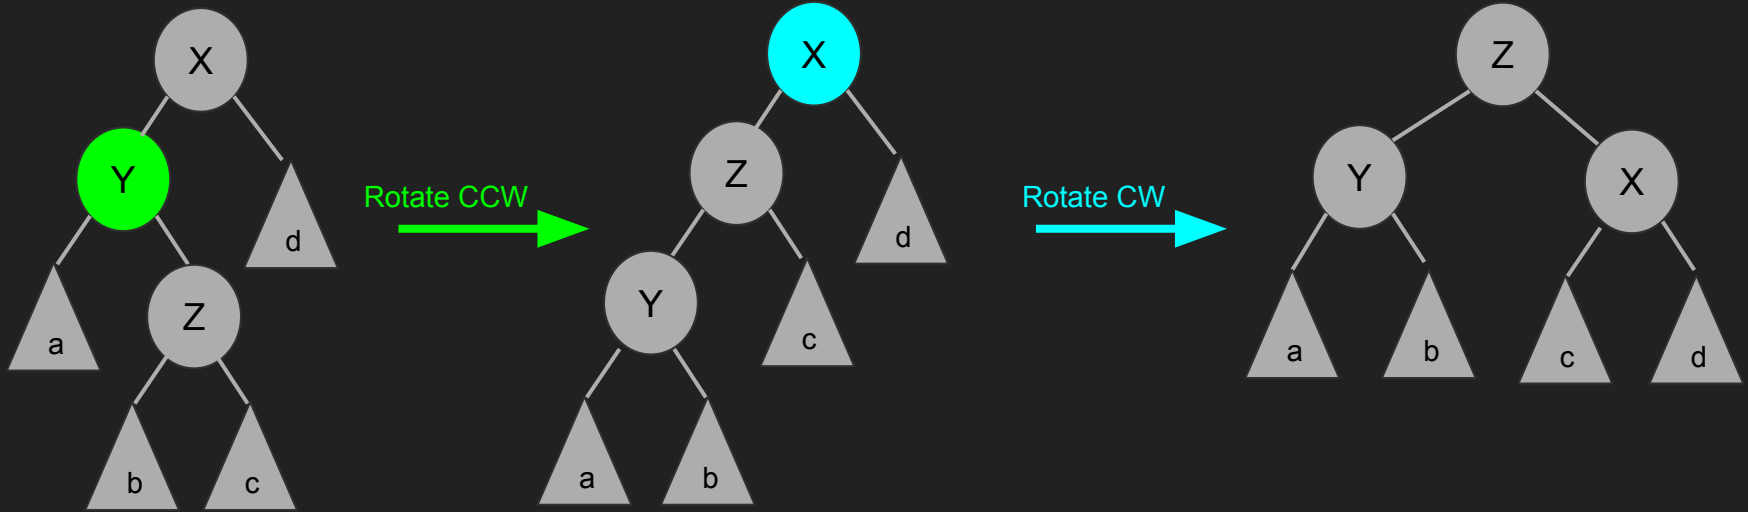
height h

AVL trees — repairing node X after Insert (subcase LR)



height h

AVL trees — repairing node X after Insert (subcase LR)




AVL trees — repairing node X after Insert (case R)


Height of right subtree of X is $h + 2$, height of left subtree of X is h

So right child of X is node Y whose height of its subtrees are h and $h + 1$

Subcase RR:




Subcase RL:



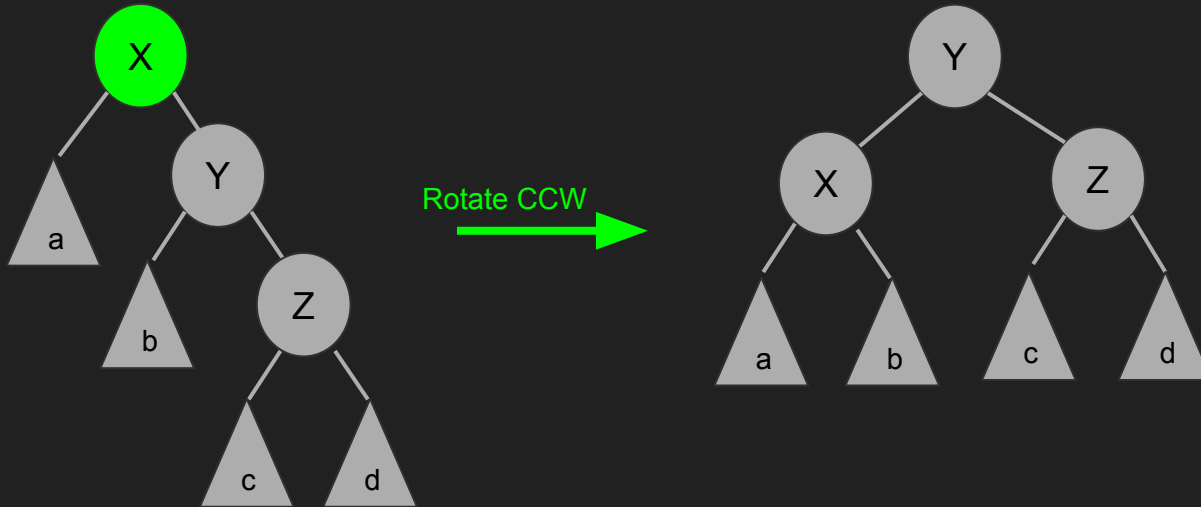
height h

AVL trees — repairing node X after Insert (subcase RR)




height h

AVL trees — repairing node X after Insert (subcase RR)



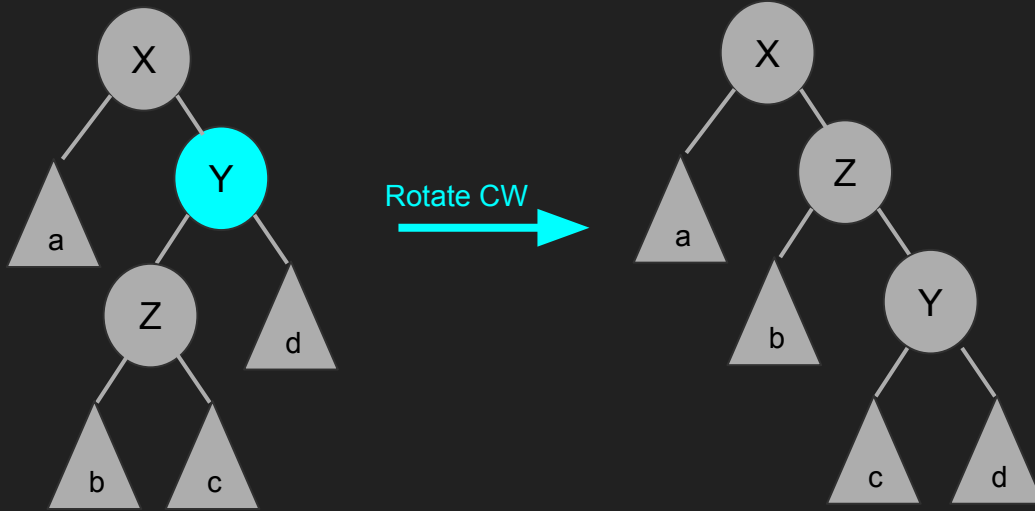
height h

AVL trees — repairing node X after Insert (subcase RL)




height h

AVL trees — repairing node X after Insert (subcase RL)



height h

AVL trees — repairing node X after Insert (subcase RL)



AVL trees — subtlety

You need to compute the height of a node efficiently

$\Theta(n)$ if you're not careful

Easiest is to store the height at each node

Update height after insert, delete, and during repairs

```
type Tree struct {  
    root *Cell  
}
```

```
type Cell struct {  
    value int  
    left *Cell  
    right *Cell  
    parent *Cell  
    height int  
}
```

Asymptotic running times

	Sorted doubly-linked list (with tail pointer)	Binary search tree	AVL tree
Search	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$
Insert	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$
Delete	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$
Minimum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$
Maximum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$

Asymptotic running times

	Sorted doubly-linked list (with tail pointer)	Binary search tree	AVL tree
Search	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$
Insert	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$
Delete	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$
Minimum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$
Maximum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$

Making Minimum and Maximum $\Theta(1)$

Store pointers to nodes holding:

- predecessor value
- successor value

at each node

Store pointers to cells with minimum value
and maximum values in tree

Update pointers during inserts and deletes

```
type Tree struct {  
    root *Cell  
    min *Cell  
    max *Cell  
}
```

```
type Cell struct {  
    value int  
    left *Cell  
    right *Cell  
    parent *Cell  
    pred *Cell  
    succ *Cell  
    height int  
}
```

Asymptotic running times

	Sorted doubly-linked list (with tail pointer)	Binary search tree	AVL tree	AVL tree (with pred / succ)
Search	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$	$\Theta(\log_2 n)$
Insert	$\Theta(n)$	$\Theta(n)$	$\Theta(\log_2 n)$	$\Theta(\log_2 n)$
Delete	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$	$\Theta(\log_2 n)$
Minimum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$	$\Theta(1)$
Maximum	$\Theta(1)$	$\Theta(n)$	$\Theta(\log_2 n)$	$\Theta(1)$