

Level 2

Optimizing Minimax

Riccardo Pucella

September 30, 2014

Summary

Round-based two-player games

- Human (or computer) versus computer

Game trees

- Records possible plays
- Decide best move from any position

Minimax algorithm

- utility of terminal states
- propagate up the game tree

Summary

Round-based two-player games

- Human (or computer) versus computer

Game trees

- Records possible plays
- Decide best move from a

Maximizing for A

Minimizing for B

Minimax algorithm

- utility of terminal states
- propagate up the game tree

Big problem: size!

Minimax works great, but it's slow

10^4 nodes/sec (not unreasonable, not great)

- 100 secs (~2 minutes): 10^6 nodes
- 1000 secs (~15 minutes): 10^7 nodes
- 10000 secs (~2.5 hours): 10^8 nodes
- ...

4x4 Tic-Tac-Toe ~ 10^{12} nodes

Four solutions

In increasing order of effectiveness

- Local optimizations
- Caching
- Pruning
- Giving up on finding **best** move

Solution 1: Local Optimizations

Pareto Principle (variant):

90% of the time is spent in 10% of the code

Find that 10%

Make sure it run fast

Optimize the deeper into the loops
(loops may be implicit via recursion)

Profiling

The best way to determine where your code is spending time is to run a profiler

For Python: `cProfile` module

(There are others. Note that a profiler has overhead. `cProfile` has low overhead.)

Example: counting primes

```
def prime (n):  
    for i in range(2,n):  
        if i > math.sqrt(n):  
            return True  
        if n % i == 0:  
            return False  
    return True
```

```
def main ():  
    count = 0  
    for i in range(20000):  
        if prime(i):  
            count += 1  
    print count
```


Example: counting primes

```
>>> import primes
```

```
>>> import cProfile
```

```
>>> cProfile.run('primes.main()')
```

```
2264
```

343540 function calls in 3.380 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	3.380	3.380	<string>:1(<module>)
1	0.013	0.013	3.380	3.380	primes.py:12(main)
20000	1.535	0.000	3.366	0.000	primes.py:4(prime)
303536	0.053	0.000	0.053	0.000	{math.sqrt}
20001	1.778	0.000	1.778	0.000	{range}

Example: counting primes

```
def prime (n):  
    if n % 2 == 0:  
        return False  
    for i in range(3,n,2):  
        if i > math.sqrt(n):  
            return True  
        if n % i == 0:  
            return False  
    return True  
  
def main ():  
    ...
```

Example: counting primes

```
>>> import primes
```

```
>>> import cProfile
```

```
>>> cProfile.run('primes.main()')
```

```
12262
```

```
177336 function calls in 0.920 seconds
```

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.920	0.920	<string>:1(<module>)
1	0.011	0.011	0.920	0.920	primes.py:14(main)
20000	0.431	0.000	0.907	0.000	primes.py:4(prime)
147332	0.025	0.000	0.025	0.000	{math.sqrt}
10001	0.452	0.000	0.452	0.000	{range}

Example: counting primes

```
def prime (n):  
    if n % 2 == 0:  
        return False  
    for i in range(3,n,2):  
        if i*i > n:  
            return True  
        if n % i == 0:  
            return False  
    return True
```

```
def main ():  
    ...
```

Example: counting primes

```
>>> import primes
```

```
>>> import cProfile
```

```
>>> cProfile.run('primes.main()')
```

```
12262
```

30004 function calls in 0.840 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.840	0.840	<string>:1(<module>)
1	0.011	0.011	0.840	0.840	primes.py:14(main)
20000	0.387	0.000	0.828	0.000	primes.py:4(prime)
10001	0.442	0.000	0.442	0.000	{range}

Example: Tic-Tac-Toe

Many functions are called at every node of the game tree:

- check if board is terminal (done)
- compute possible moves
- apply a move to the board

Let's analyze `has_win` (part of `done`)

Function has_win

board is array of 0, 1 (for 0), 10 (for X)

```
def has_win (board):  
    for positions in WIN_SEQUENCES:  
        s = sum(board[pos] for pos in positions)  
        if s == 3:  
            return '0'  
        if s == 30:  
            return 'X'  
    return False
```

```
WIN_SEQUENCES = [  
    [0,1,2],  
    [3,4,5],  
    [6,7,8],  
    [0,3,6],  
    [1,4,7],  
    [2,5,8],  
    [0,4,8],  
    [2,4,6]  
]
```

Function has_win

board is array of 0, 1 (for 0), 10 (for X)

done

3228033 function calls in 1.826 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	1.826	1.826	<string>:1(<module>)
100000	0.590	0.000	1.766	0.000	done.py:166(has_win_2)
2502424	0.561	0.000	0.561	0.000	done.py:168(<genexpr>)
1	0.060	0.060	1.826	1.826	done.py:27(test)
625606	0.615	0.000	1.176	0.000	{sum}

Function has_win

board is array of 0, 1 (for 0), 10 (for X)

checks only rows that involve last move

```
def has_win ((board,last_move)):
```

```
    for positions in CHECKS[last_move]:
```

```
        s = sum(board[pos] for pos in positions)
```

```
        if s == 3:
```

```
            return '0'
```

```
        if s == 30:
```

```
            return 'X'
```

```
    return False
```

```
CHECKS = {
```

```
    0: [[1,2],[3,6],[4,8]],
```

```
    1: [[0,2],[4,7]],
```

```
    2: [[0,1],[5,8],[4,6]],
```

```
    3: [[4,5],[0,6]],
```

```
    4: [[1,7],[3,5],[0,8],[2,6]],
```

```
    5: [[3,4],[2,8]],
```

```
    6: [[7,8],[0,3],[2,4]],
```

```
    7: [[6,8],[1,4]],
```

```
    8: [[6,7],[2,5],[0,4]]
```

```
}
```

Function has_win

```
# board is array of 0, 1 (for O), 10 (for X)
# checks only rows that involve last move
```

1164907 function calls in 0.789 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.789	0.789	<string>:1(<module>)
100000	0.307	0.000	0.730	0.000	done.py:175(has_win_3)
798678	0.199	0.000	0.199	0.000	done.py:177(<genexpr>)
1	0.058	0.058	0.789	0.789	done.py:27(test)
266226	0.225	0.000	0.424	0.000	{sum}

Function has_win

board is array of 0, 1 (for 0), 10 (for X)

```
def has_win_0 (board):
    if (board[0] == board[1] and board[0] == board[2]):
        return board[0]
    if (board[3] == board[4] and board[3] == board[5]):
        return board[3]
    if (board[6] == board[7] and board[6] == board[8]):
        return board[6]
    if (board[0] == board[3] and board[0] == board[6]):
        return board[0]
    if (board[1] == board[4] and board[1] == board[7]):
        return board[1]
    if (board[2] == board[5] and board[2] == board[8]):
        return board[2]
    if (board[0] == board[4] and board[0] == board[8]):
        return board[0]
    if (board[2] == board[4] and board[2] == board[6]):
        return board[2]
    return False
```

Function has_win

board is array of 0, 1 (for O), 10 (for X)

def

100003 function calls in 0.158 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.158	0.158	<string>:1(<module>)
1	0.035	0.035	0.158	0.158	done.py:27(test)
100000	0.123	0.000	0.123	0.000	done.py:80(has_win_0)

return False

About optimizations

Effectiveness of an optimization depends on

- algorithmic choices
- data representation choices
- programming language choices
- details of the implementation of the language

In Python, copying is expensive, update is not

- In other languages, update is expensive, copying is not

Solution 2: Caching

Minimax will compute the minimax value of a board every time it encounters it during traversal of the game tree

The same board may appear as the result of a different sequence of moves

E.g. moves 0 (for X), 1 (for O), 2 (for X)

and 2 (for X), 1 (for O), 0 (for X)

both produce the same board

Basics of caching

Caching (or memoization) is an approach to remembering previous results of a function.

```
def foo (x):
```

```
    code for foo(x)
```

```
    return v
```

Basics of caching

Caching (or memoization) is an approach to remembering previous results of a function.

```
def foo (x):
```

```
    if x has been seen before, return value[x]
```

```
    code for foo(x)
```

```
    save value[x] = v
```

```
    return v
```


Basics of caching

Caching (or memoization) is an approach to remembering previous results of a function.

```
def foo (x):
```

```
    if x has been seen before, return value[x]
```

```
    code for foo(x)
```

```
    save value[x] = v
```

```
    return v
```

Lookups should be reasonably fast

They get performed at every node of the game tree

Basics of caching

Caching (or memoization) is an approach to remembering previous results of a function.

```
def foo (x):
```

```
    if x has been seen before, return value[x]
```

```
    code for foo(x)
```

```
    save value[x] = v
```

```
    return v
```

Data structure:
hash tables (dictionaries)

Lookups and saves:
constant time (mostly)

Minimax application

Keep a dictionary associating with seen boards their computed minimax value

Shortcut minimax computation when a board has been seen

Technical problem: arrays cannot be used as keys in Python dictionaries

- Need a way to associate with every board a key by which to refer to it in the dictionary

Caching problems

The caching table can get large

If the caching table gets too large, it can overflow memory, and then the system starts swapping to disk

- Any time advantage is lost because memory swapping is SUPER SLOW

Exploiting symmetry

The caching table can be used to remember more than just seen boards

It can be used to remember boards that have not been seen yet

Observation: two symmetrical boards must have the same minimax value

(Why?)

Exploiting symmetry

Two approaches:

1. When saving a minimax value for a board, save that minimax value for all symmetric boards
(space expensive)
2. When looking up a board in the table, also look for any symmetry of that board
(time expensive)


Solution 3: Pruning

When doing minimax, at a maximum node, you compute (recursively) the minimax value of the children.


If at any point you can tell that the best value you will get for a child is less than your current max, you can get stop computing the minimax value of that child (prune the subtree)

Similarly when at a minimum node


α - β pruning example




α - β pruning example




α - β pruning example



α - β pruning example



α - β pruning example



Properties of α - β pruning

Returns the same result as standard minimax

Effectiveness depends on move ordering

- Best case: can double the search depth

Can be tricky to get right in the presence of caching (why?)

The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

Something similar for Min-Value

Solution 4: Cutting-off search

For some games, none of the above is enough to make minimax manageable.

Chess, Go, Ultimate Tic-Tac-Toe.

Give up on searching the whole game tree

- Search only to a limited depth
- Possibly some branches deeper than others

What's the issue?

Evaluating nodes

Minimax works by propagating up the utility of final states

- Utility = is final state a win or not?

If you don't search to the final states, what do you propagate up?

Evaluation function:

- Associate a value that tries to capture how good the position is when cutting off

Evaluation functions

For chess, typically linear weighted sum of features

$$\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

E.g., $w_1 = 9$ with

$f_1(s) = (\# \text{ white queens}) - (\# \text{ black queens})$
etc...

Minimax with cutoff

Minimax/cutoff is similar to Minimax:

- Terminal? is replaced by Cutoff?
- Utility is replaced by Eval

In practice, for chess, searching 10^6 moves:

$$br^{\text{depth}} = 10^6, br = 35 \rightarrow \text{depth} = 4$$

4-ply lookahead is a hopeless chess player!

- 4-ply \approx human novice
- 8-ply \approx typical PC, human master
- 12-ply \approx Deep Blue, Kasparov